

Lecture #1: Overview & Design Principles #1

WPI CS4516 Spring 2019 D term

*Instructor: Lorenzo De Carli (ldecarli@wpi.edu)
(slides include material from Christos Papadopoulos, CSU)*



Course topics

- Review of **foundational topic** in computer networks
- Based on **reading list of scientific papers from networking literature**
 - You are responsible for reading each paper **before lecture**
- Focus on works with **long-lasting, practical impact** on networking
 - If interested in **cutting-edge research topics** and/or **theoretical work** look at **graduate-level offerings**



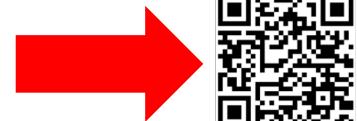
Course organization

- Every week we will cover a relevant networking topic
 - Internet design principles
 - Network layer
 - Transport layer
 - Network measurements and monitoring
 - Network management (2 weeks)
 - Router design
 - Net neutrality
- During every lecture we will discuss an **influential paper** related to that week's topic



Course material

- **Reading list** of 13 papers
 - **No textbook**
 - I may provide **optional material** to aid understanding
- **Lecture slides**
- You are expected to **read each paper before coming to lecture**
- **Discussion** with your classmates is **encouraged**
- I will set up **forums on Canvas** that you can use to **engage in conversation** with your classmates and the instructor



Resources & course management

- Course overview, syllabus, assigned readings on class website:
 - <https://web.cs.wpi.edu/~ldecarli/teaching/cs4516D19>
- Same material also available on Canvas
 - <https://canvas.wpi.edu/courses/12976>
- Canvas will be used for uploading assignments, taking quizzes, interacting w/ students & instructor, posting grades



Class project

- The class also requires to complete a **practical project**
- Project goal: implement a gateway to perform **application identification** from **encrypted traffic**
- **4 project deliverables** to be completed at different times during the term
- The project will be introduced in the next lecture



Grading

- Six 20-minute **in-class quizzes**: 60% of grade
 - Each quiz at the **beginning of each week**, except for week #1 (**no quiz**) and week #7 (**quiz at the end of the week**)
 - Each quiz covers the **material from the previous week** (material discussed in class and papers)
 - Just to be clear: **content of the papers** (even if not discussed in class) **is fair game for quizzes**
 - I will **drop the lowest quiz score** when computing the final grade
 - This also mean **you can miss one quiz w/o impacting your grades**



Grading - II

- Class project: 40% of grade
 - Project will consist of **four phases**
 - At the end of phases 1-3, you will need to deliver **what you implemented so far** (each phase builds on the previous one)
 - At the end of phase 4, you will need to deliver a **final report and the project implementation**
 - Each phase will receive a separate grade. The overall quality of the implementation will also be evaluated.



Participation, behavior etc.

- “Good behavior coefficient” (GBC)
- Final grade:
 - $(0.6 * \text{quiz score} + 0.4 * \text{project score}) * \text{GBC}$
- Takes into account **attendance** and **respectful behavior towards instructor, TAs, classmates** (or lack thereof)
- Can also go > 1.0 ☺
- (If you receive a GBC $\neq 1$, I will tell you why)
- (Will also try to warn you, as reasonable, if **your behavior is at risk of affecting your grade**)



Lastly, who am I?

- **Lorenzo De Carli**
 - **Assistant Professor of Computer Science**
 - Politecnico di Torino (Torino, Italy):
 - B.Sc. (2005), M.Sc. (2007) in Computer Engineering
 - University of Wisconsin-Madison:
 - M.Sc. (2010), Ph.D. (2016) In Computer Science
 - Assistant Professor of CS at Colorado State University from 2017 to 2018
 - Now at WPI ☺
 - Interested in **Networking, Programming Languages, Security**

One more thing before we begin...



How to read papers

- In order to successfully understand scientific papers, it is important to learn **how** to read them
- **Common mistake:** read a paper sequentially the way you would read a book, dedicating equal attention to all parts
 - **It almost never works** - you will drown in the details, miss the big picture, and get very frustrated
 - Instead, try reading a paper selectively, **focusing on the most important parts first**



How to read papers - II

- The trick is to read a paper **multiple times**:
 - **First step:** read abstract, intro and conclusion; skim the beginning of every section in between. **Goal:** understand which problem the paper aims to solve, and get a very basic understanding of how it gets solved
 - **Second step:** read the paper end-to-end, trying to understand the technical details of how problems are solved.
 - For particularly complex papers, **you may need to repeat the second step multiple times**, every time delving deeper into the technical details
- **Additional readings** (linked on the course web page):
 - M. Hanson, “Efficient Reading of Papers in Science and Technology”
 - S. Keshav, “How to Read a Paper”



How to read papers - III

- **Summarizing:** skim the paper first to understand what is the problem being solved, and then try to refine your understanding with further readings
- What happens if you follow this method:
 - **Pass 1:** *this paper proposes some hardware which can execute different networking algorithms*
 - **Pass 2:** *this paper propose a dataflow-based processor which can efficiently run multiple forwarding and classification algorithms, overcoming the limitations of traditional network ASICs*



How to read papers - IV

- What happens if you don't follow this method and try to understand everything at once:
 - Pass 1: this paper solves a problem in networking
 - Pass 2: this paper solves a problem in networking
 - Pass 3: this paper solves a problem in networking
 - Pass 4: this paper solves a problem
 - Pass 5: this paper
 - Pass 6: ZZzzzz...



How to read papers - V

- Other tips:
 - **Take notes**
 - Form informal **reading groups** with your classmates so you can read and discuss the paper together
 - **Come to the lecture prepared for discussion** (especially if there are aspects of a paper which you did not understand)

Now, let's get down to
business...

Internet design #1: how to design complex systems

Reading: “End-to-end Arguments in System Design”

J.H. Saltzer, D.P Reed & D.D. Clark



Let's begin from the beginning

- Systems of networked computers (e.g., the Internet) are **incredibly complex objects**
 - **Too complex for us humans** to fully grasp!
 - So how do us, puny humans, go about building one?
- We need a plan!
- ... or more precisely, some **guiding principles**...
- ...on which **everyone** agrees



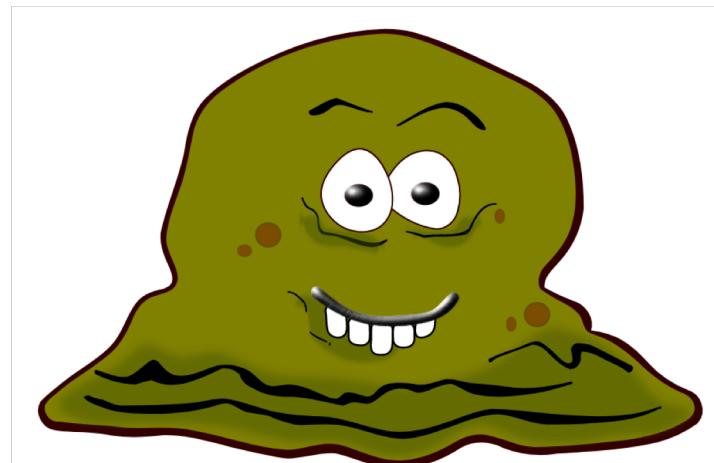
A general principle

- “**Modularity guided by abstraction** is the way **things get done**” (Scott Shenker)
 - Can you explain me what this means?
- Pretty much **every idea that everyone has ever had** in the field of designing computer systems implements this principle!

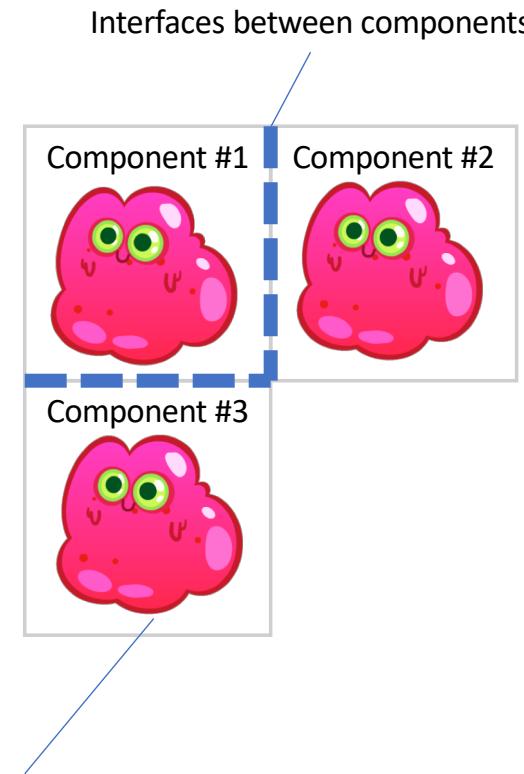
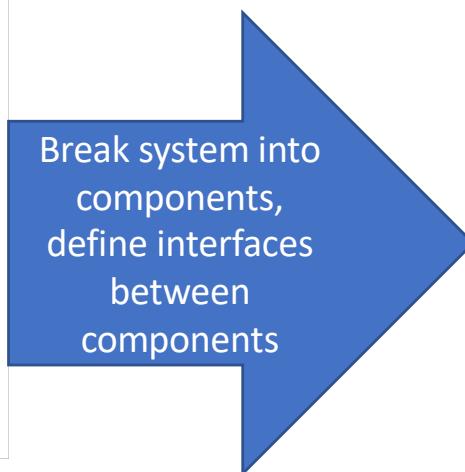


So, this is what we need to do...

Suppose that we need to design a system with **horrendous complexity**...



Horrendous complexity



Manageable complexity



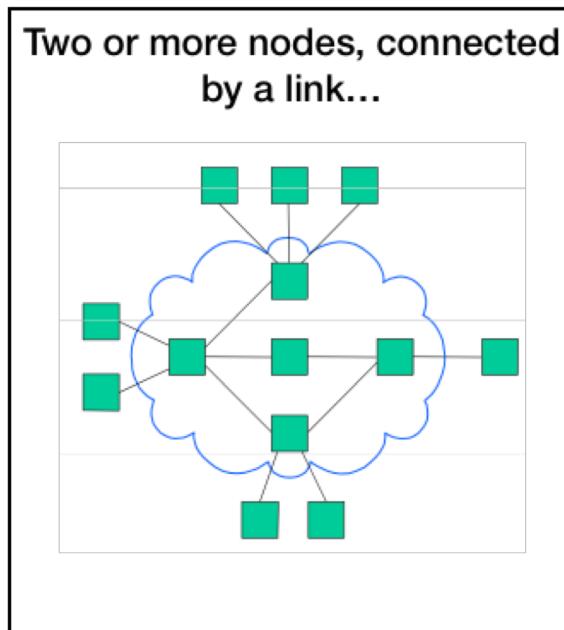
This is still too general

- The idea of breaking a complex system into simpler components is powerful...
- ...but still too general!
- How to translate it to concrete design principles?
 - Particularly in the domain of **designing a network?**

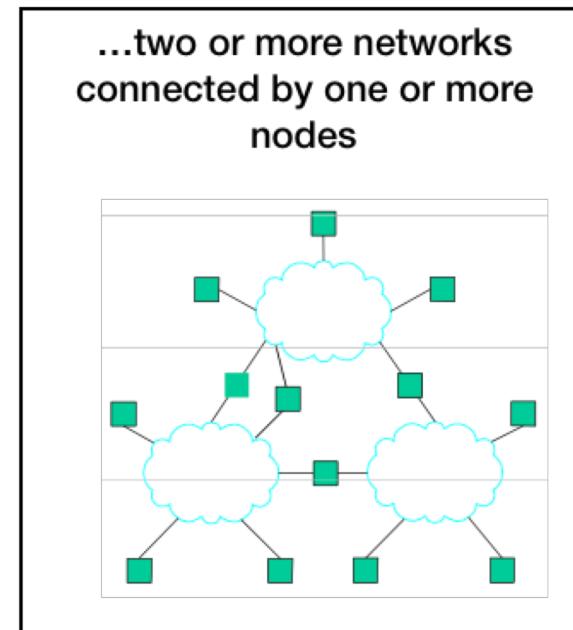


First things first

- What is a network?
 - A computer network can be recursively defined as:



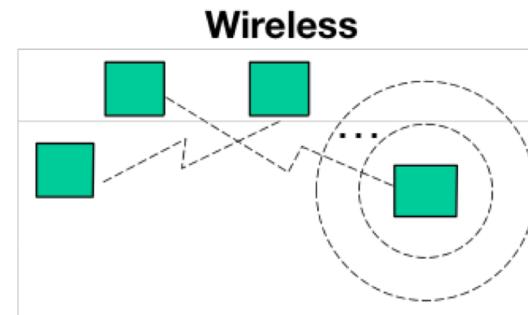
...OR...





Nodes and links

- A **node** is pretty much any computing device that can communicate with other devices (PC, phone, smart device, car, etc.)
- A **link** is a component (physical or logical) which two or more nodes use to exchange data
- **Examples of links:**





Where is the complexity in a networked system?

- Any guesses?
- Some core problems:
 - Making effective use of the physical medium (relatively easy w/ wires, less so with wireless)
 - Figuring out how to reach a specific destination
 - Ensuring communication can continue even if individual links fail
 - Multiplex many communication flows on the same links
 - Allowing different applications to communicate using the same network infrastructure
 - ... (more on this in lecture #2)



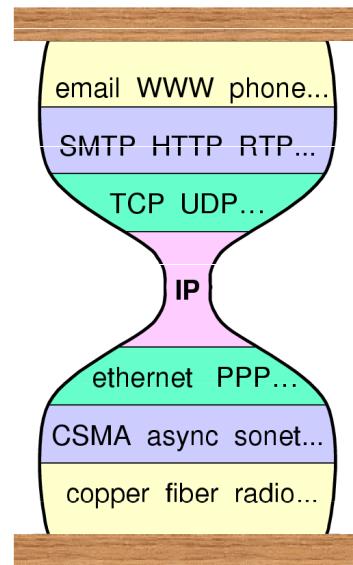
How do networks break this complexity?

1. Identify a **set of core tasks** that the network infrastructure must perform
2. Define **how** the infrastructure must perform each task
 1. These definitions are usually called **protocols**
3. Define how protocols **interact** with each other



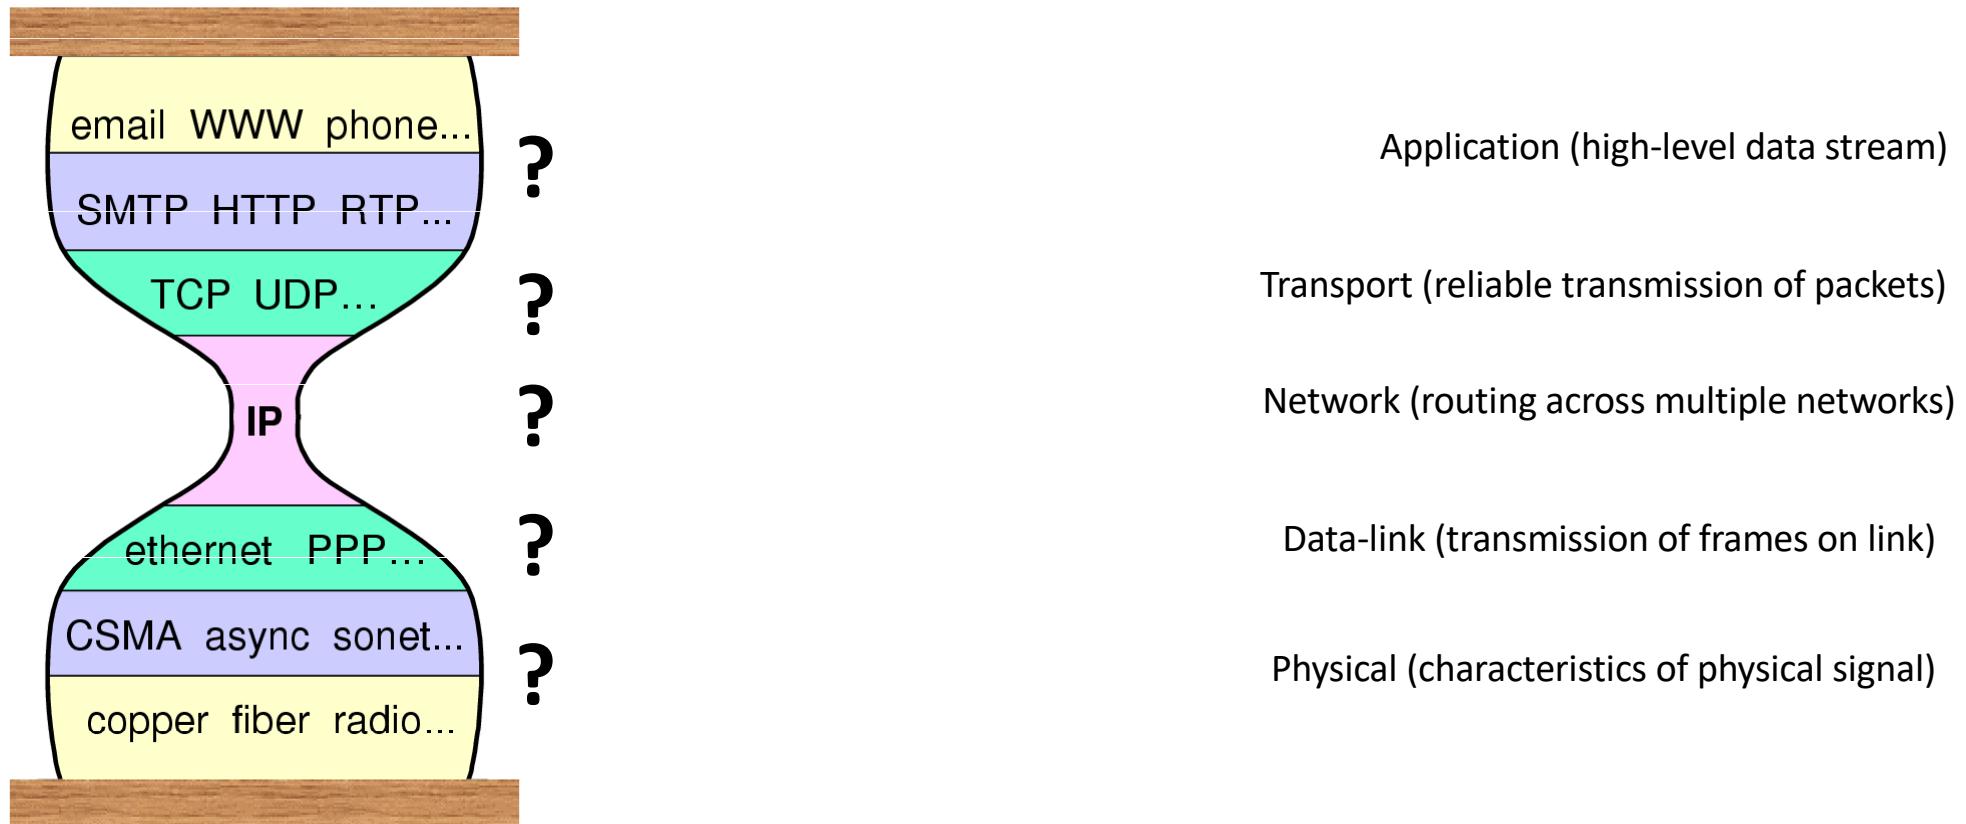
Protocol stack

- Computer networks are based on a **stack of protocols** which provide features at **increasingly high level of abstraction**
- The **hourglass design**:





Review time!



Protocol stack - II

- Today we won't discuss in the details the functions of each layer...
- Instead, we will review a **general principle** that helps designers to decide **how to distribute functions across layers**



End-to-end arguments in system design

- 1981 paper by Jerry Saltzer, David Reed, Dave Clark
 - MIT Laboratory for Computer Science
 - All authors went on to make important contribution to the Internet design and architecture
- **Heuristic** for solving the problem of **function placement** among the modules of a distributed system
- Highly influential



The end-to-end argument

- Architectural principle
- Deals with where to place protocol functionality
 - Inside the network (i.e. in switches/routers), or
 - At the edges (endpoints)



The core principle

If a function can only be fully and correctly implemented with the help of the applications standing at the endpoints of a communication system, providing that function as a feature of the communication system is not advisable



Can you explain what it means?



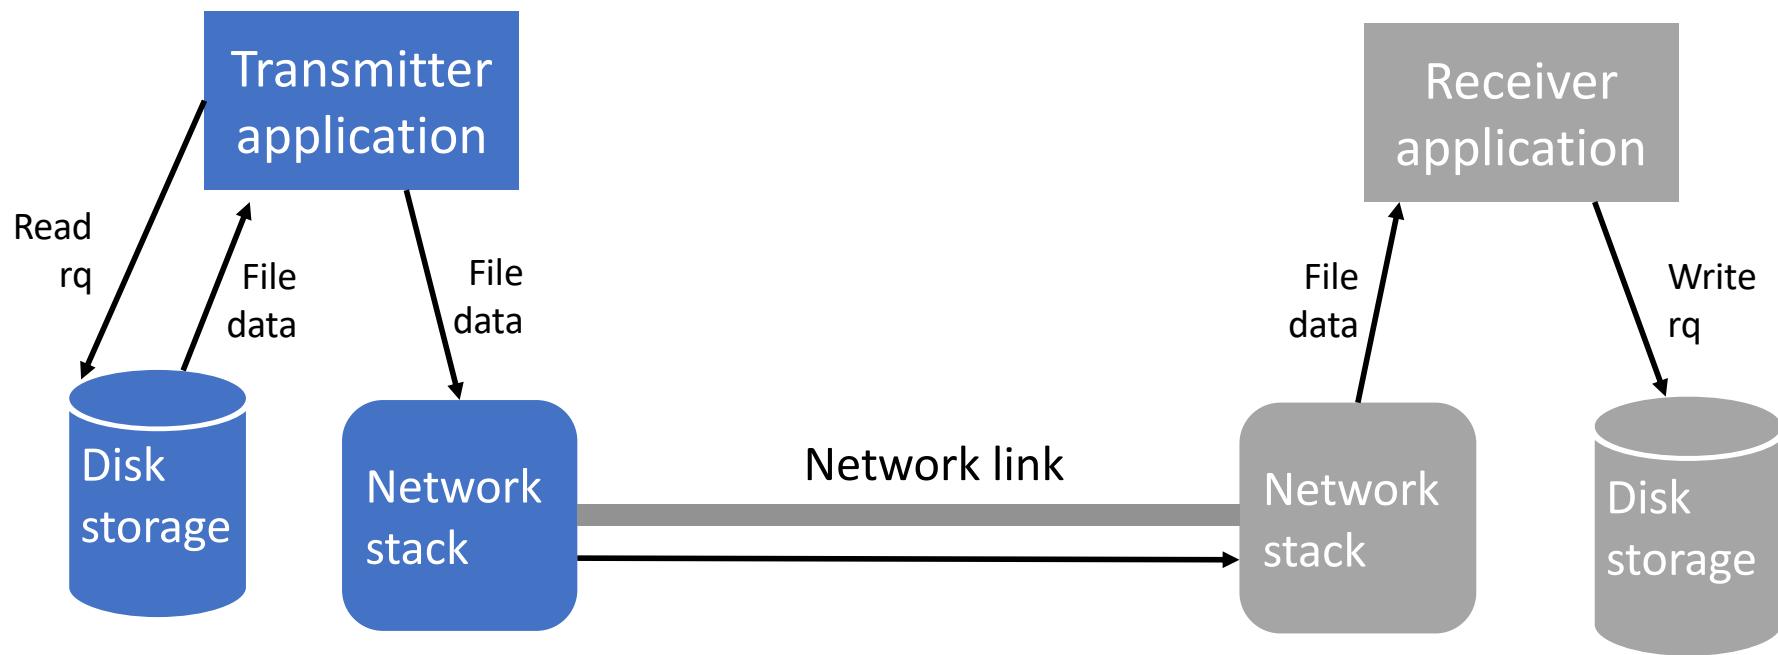
Motivation

- Don't duplicate functionality at multiple levels if you have to do it at the top anyway
- Why? Two main reasons:
 - Duplicate functionality has a cost associate with it (e.g. reliable transport introduces delays and retransmissions)
 - Functionality provided at low level is forced upon all applications, even those which do not need it or would work better without it



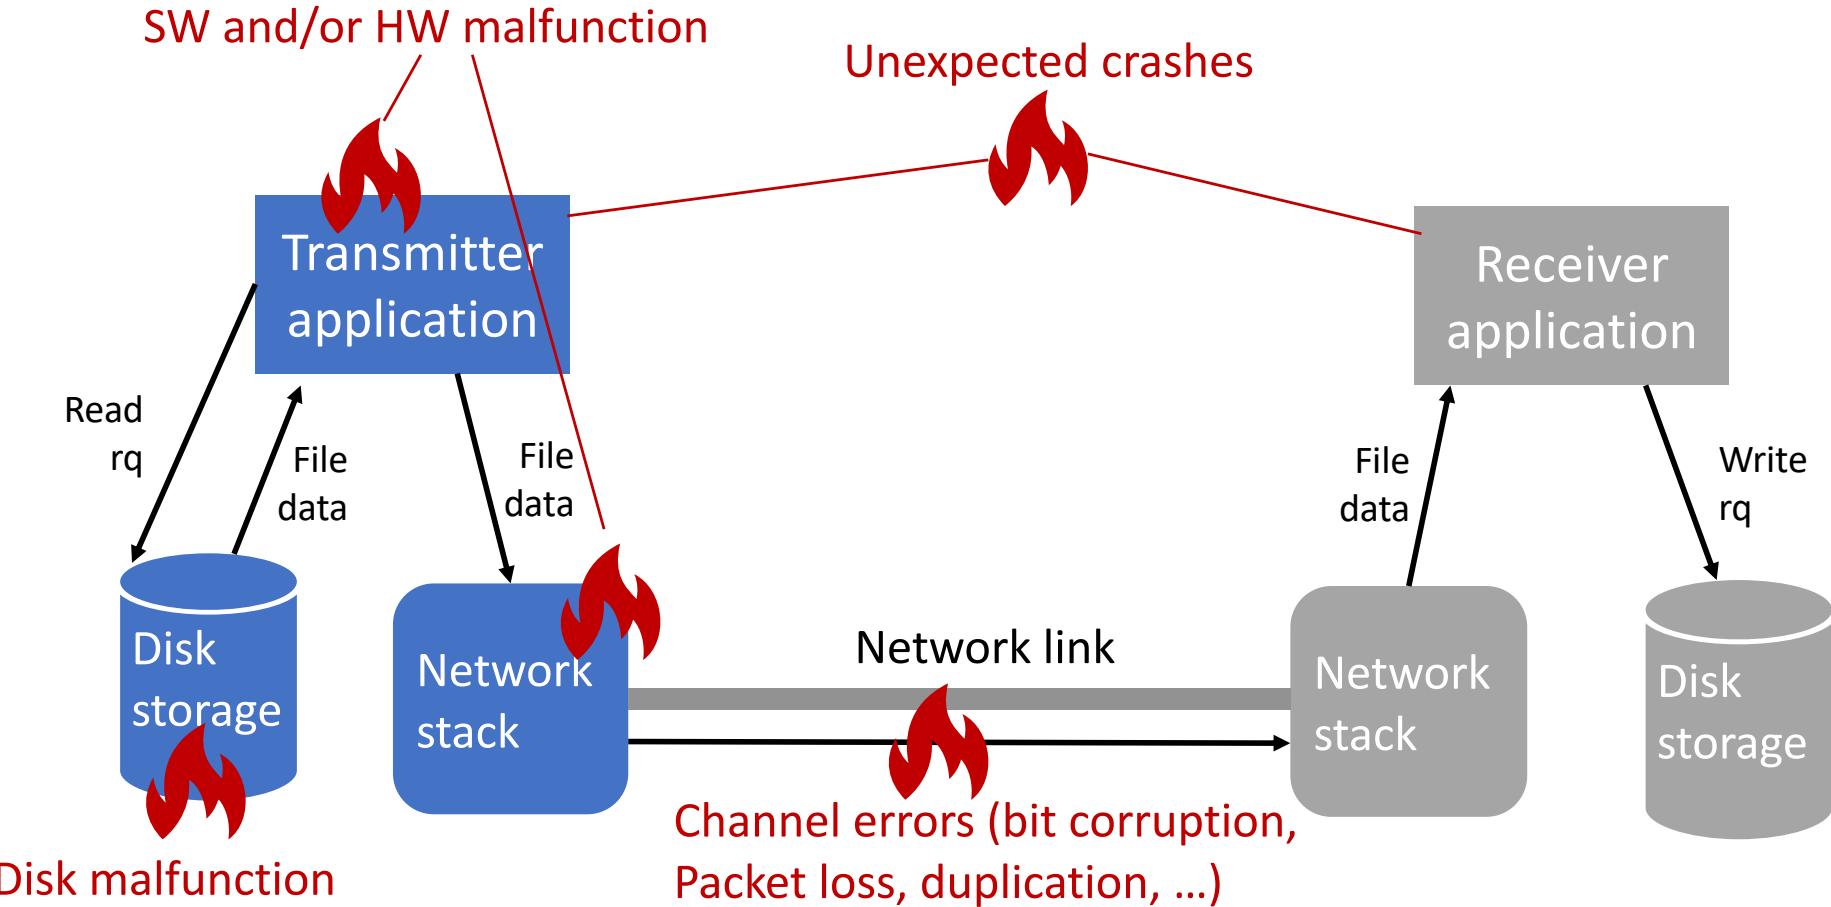
Example

- **File transfer application**
- Components:



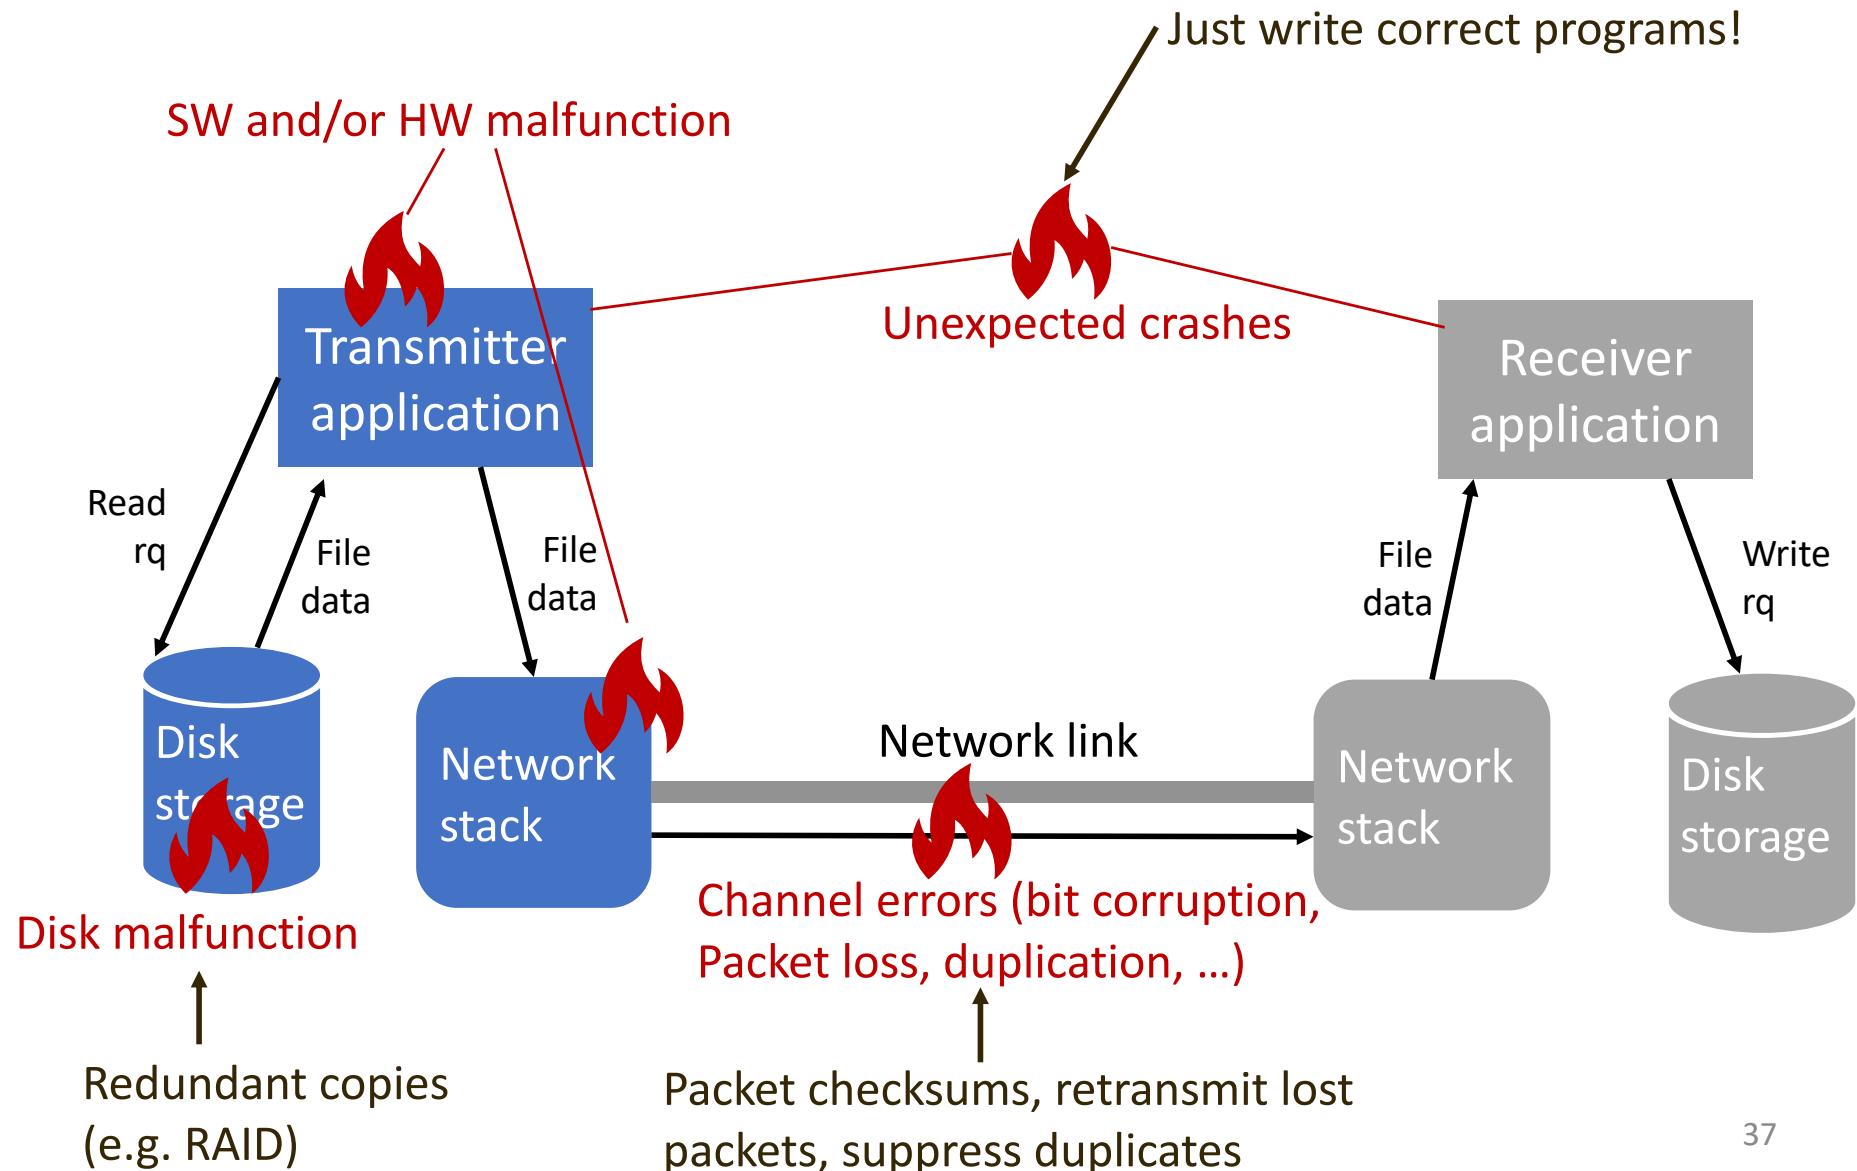


What could possibly go wrong?



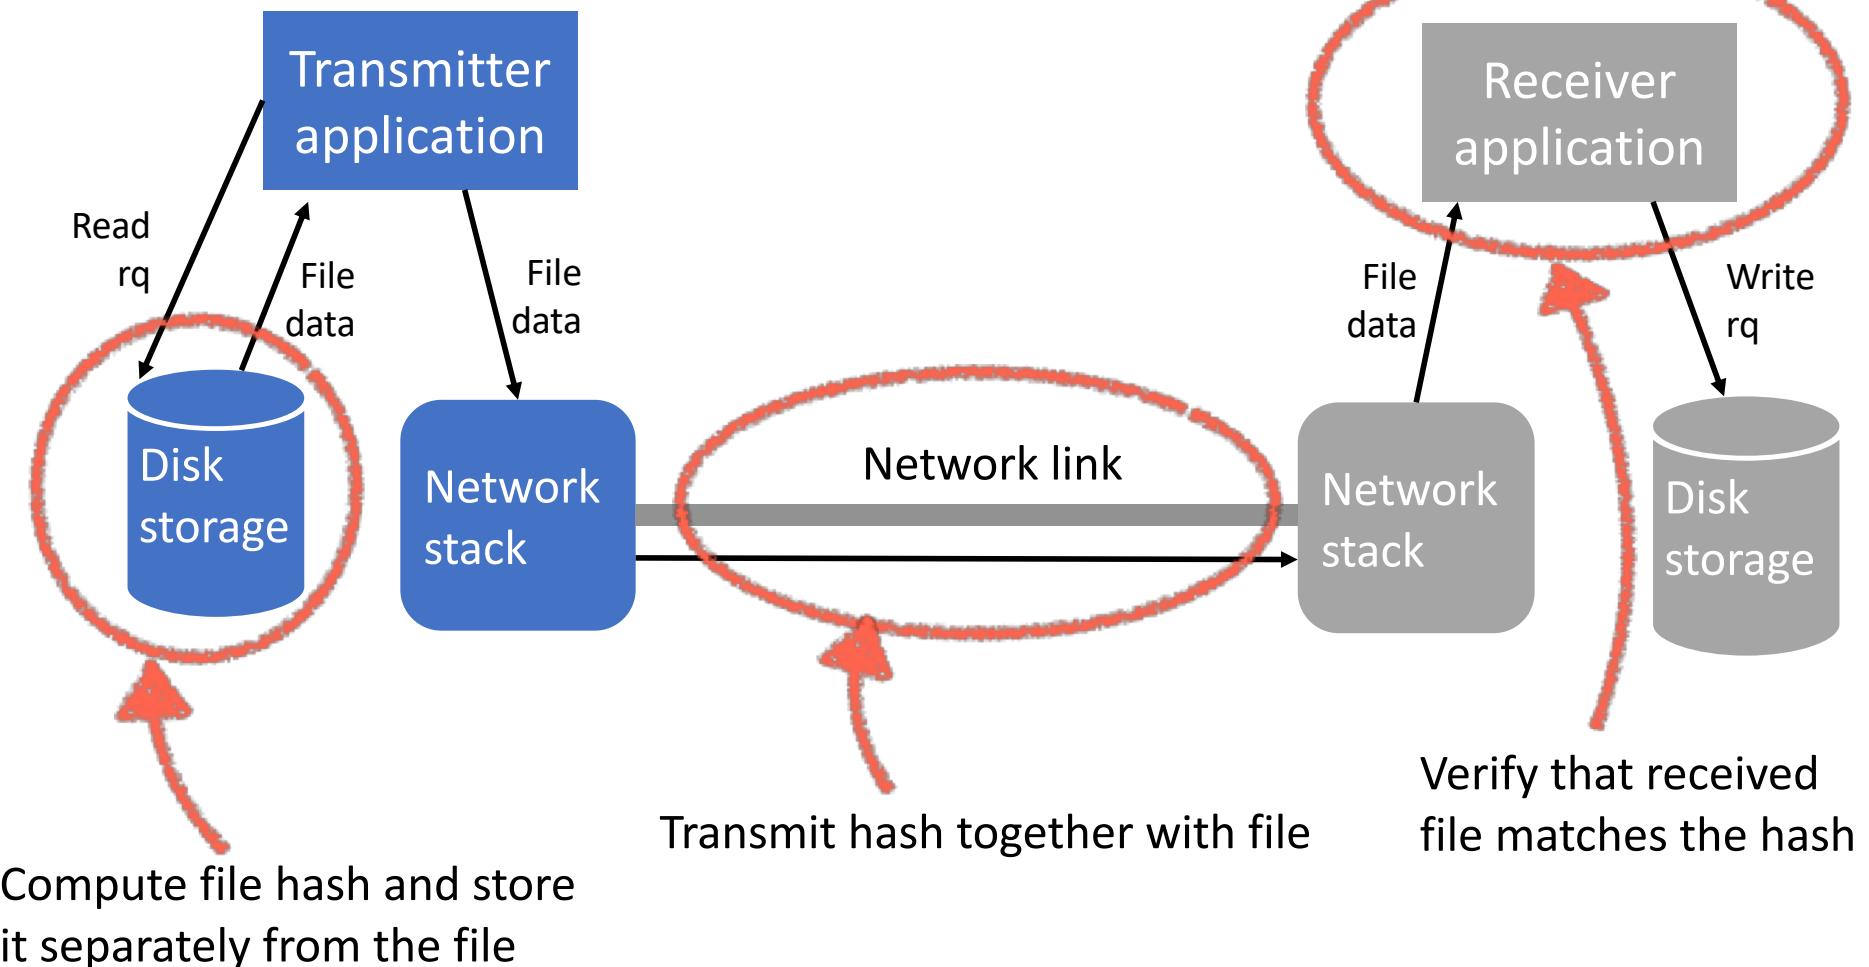


How would you go about solving the problem?





You could do all of this... or ...



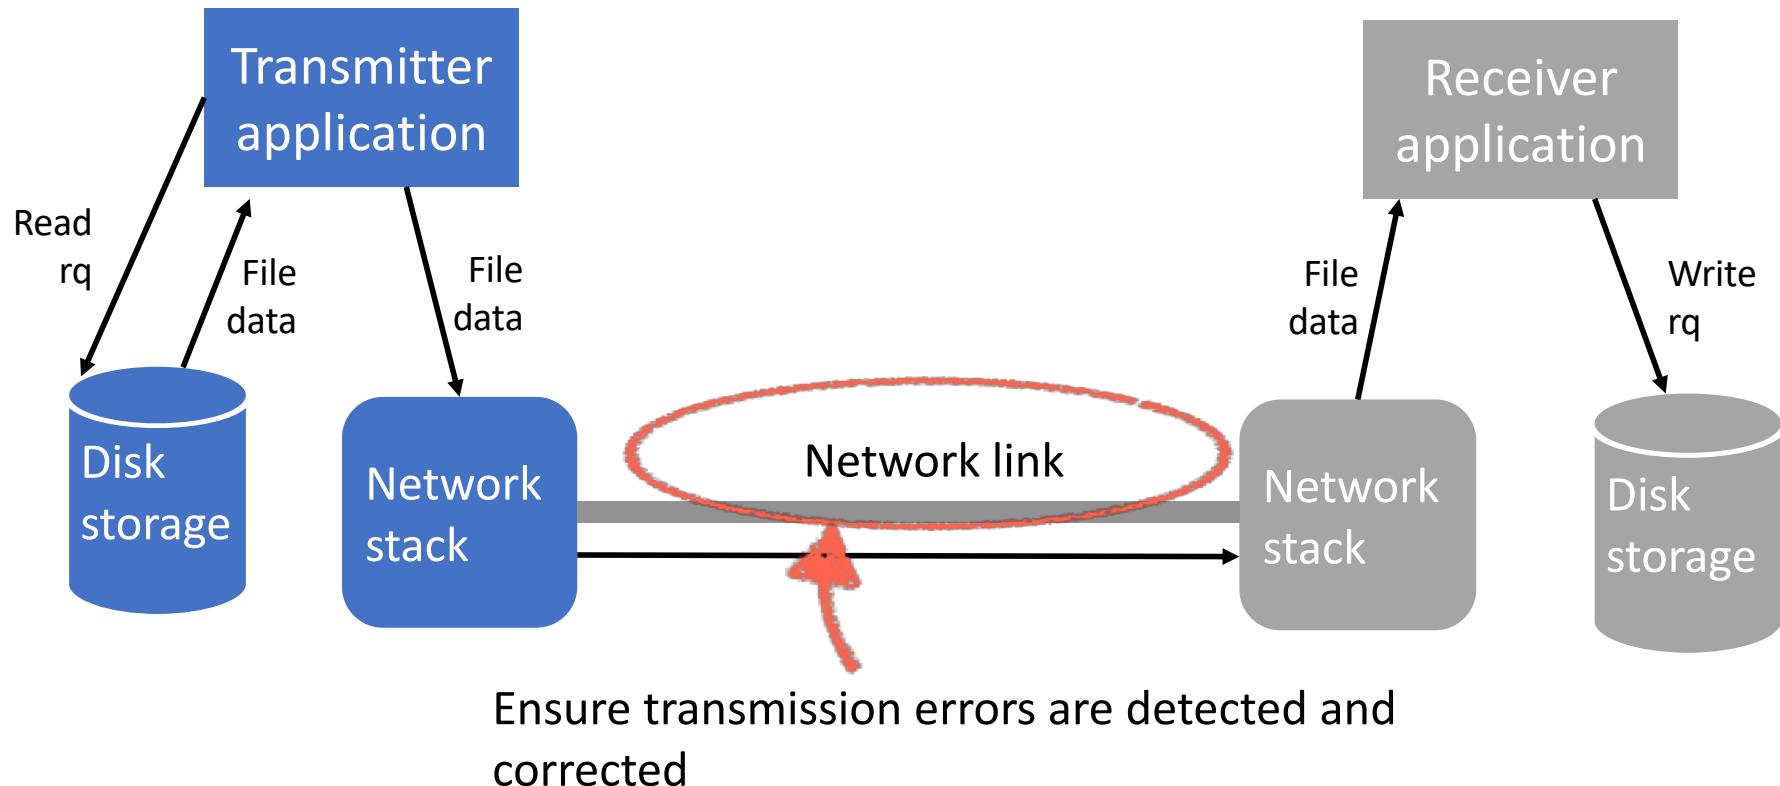


Why is this a good idea?

- Guards against all malfunctions we have considered: errors are discovered regardless if they originate in faulty disks, SW/HW or network links
- Simple and economical: requires no built-in redundancy and error correction in the system
- If everything works correctly, most transfers will be completed at first attempt
 - If a transfer require more than a couple of tries to be completed, there is probably an issue in the system



Alternative approach: reliable network transmission



Does not solve the problem: I still need to transmit and check the file hash to guard against all other possible errors



Another drawback of not respecting e2e

- False sense of security
- Why?
- MIT code repository example:



- Result: OS source gradually corrupted!



Everything is relative

- Correctness checks in individual components do not solve the reliability issue...
 - But they may make sense as optimizations!
- Consider a file which is transmitted over **100** packets. Probability of packet error is p_E . What is the expected value $E[T]$ of the number of trials T before the file is transmitted successfully, for various values of p_E ?

p	$E[T]$	
0.00001	~1	
0.0001	~1	
0.001	1.10	
0.01	2.73	
0.1	37648	

} Acceptable
} Maybe acceptable?
} Definitely not acceptable!

What's the lesson here?

- Violations of e2e, while not recommended in general, may make sense as **optimizations**
- In this example, a highly unreliable network may reduce the chance of successful transmission to negligible values
- It makes sense to introduce reliability as a network feature (e.g. automatically retransmit corrupted packets)
- When is this acceptable? As always in engineering, the answer is “it depends” :-)

Some case studies

Delivery guarantees

- Should the network acknowledge message reception?
 - **Argument in favor:** increases reliability, message integrity etc.
 - **Argument against:** increases delay - network stack must wait for acknowledgement before further transmission
 - What if app does not need it?
 - Can you think of an example?

Delivery guarantees - II

- The argument against delivery guarantees in the paper is that the sender really cares about the **effect** of the packet, not its **reception**
 - Distributed protocols, databases, etc. all incorporate some sort of explicit application-level commit - having a network-level acknowledgement is not enough!

Delivery guarantees - III

- The argument against delivery guarantees in the paper is that the sender really cares about the **effect** of the packet, not its **reception**
- **Can you think why this consideration may be outdated in some cases?**

Secure transmission of data

- Can you name a technology for true end-to-end data encryption?
 - IPsec - end-to-end-ish: encryption between IP endpoints, but not between applications
 - TLS?

Secure transmission - II

- Why is it a good idea to have e2e encryption?



- Why is it a good idea to have non-e2e encryption?

Guarantees encryption of all traffic!

Ordered message delivery

- May be difficult or impossible to provide within the network
 - E.g. packets may take different paths in the network due to transient routing failures
 - In that case, packets may arrive out of order - and there is little any entity participating in the communication can do

Duplicate message suppression

- Can you think of a reason why doing this in the network is not enough?

Do not refresh the page while your payment is being processed

Processing may take a few minutes if our system is very busy

If you are unsure or you don't get a confirmation screen after processing please ring the following number in normal office hours **01708 434343**

Alternatively you can email us on ConfirmationPayment@havering.gov.uk

[Cancel payment and go back](#)

[Previous page](#)

[Process now](#)



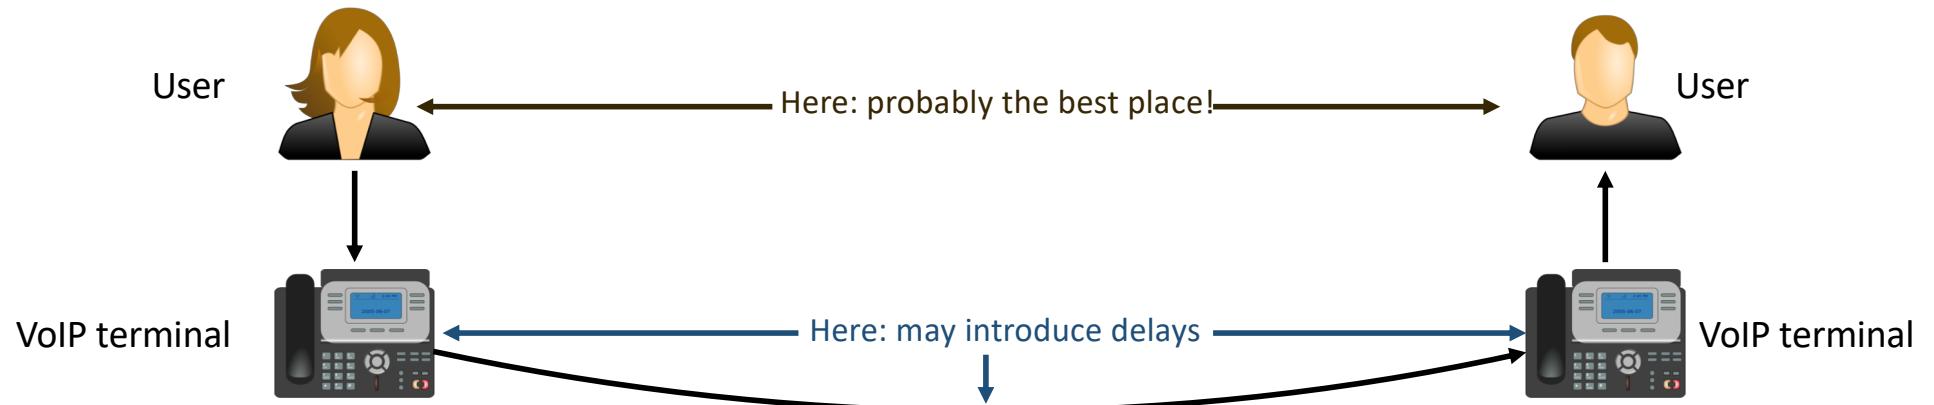
! You can always go back through the pages by pressing the **Previous Page** button if you need to change some of the details

Nuances

- E2e specifies that functionality should be implemented at the endpoints - but it does not specified what those endpoints are
- In certain cases, the choice of endpoints is non-obvious and has important consequences on the system design

What is an endpoint?

- Example: voice communication



- Where should error correction go?

What is an endpoint? - II

- This use case does not imply that you should always let the user take care of the problem :-)
- Consider the case of a voicemail server which streams pre-recorded voice messages
 - Error correction is probably a good idea there
- **Lesson:** e2e cannot be applied blindly - designers must use some care in choosing the correct endpoints!

Not just for networking!

- E2e is often presented as a principle for distributed system design
- Which is interesting, because the paper presents several examples not related to distributed systems
- Do you remember which ones?

CPU design

- The old CISC vs RISC debate
 - CISC: complex functionality provided as processor instructions
 - RISC: processor only provides basic operation (conditional, arithmetic, etc.) - everything else build in software

OS design

- Monolithic OS vs microkernel
 - Monolithic: OS implements an extensive amount of functionality (networking, file system, etc.) in kernel space
 - Microkernel: OS provides isolation and interprocess communication, everything else is implemented outside the kernel (and does not affect its stability)
- Which one best implements the e2e idea?

Can we reformulate e2e in a non-networking specific way?

- *The correct level of abstraction at which to provide a function in a complex system is the level at which the most information about the function's purpose is available*
 - **CPU design:** complex instructions should be implemented in software, not in hardware
 - **OS design:** all primitives but process isolation and IPC should be implemented by user space libraries, not the kernel
 - **Compiler design:** program optimization should be performed by the compiler, not the programmer (“premature optimization is the root of all evil”, Knuth)
 - **Network design:** high-level functionality should be implemented at the endpoints, not in the network core

A final consideration

- E2e is not a law - it is a **heuristic**
 - *Heuristic: a technique serving as an aid to problem solving by experimental and especially trial-and-error method* (Merriam-Webster dictionary)
- The paper we just discussed condenses lessons from years of experience in system design - but by no means this mean that this is always the correct approach
- E2e fosters **simplicity** and **elegance** of design, which are good things...
- But practical systems are often neither elegant nor simple!

Occam's razor

- Heuristic principle to help in scientific investigation (and any investigation, really)
- Often misquoted as “the simplest solution is most likely correct”
- It should actually be interpreted as “the simplest solution **among all the plausible ones** is most likely correct”
- Keep in mind that **correctness > simplicity**

Breaking e2e

- Can you think of networking technologies that break e2e and yet are widely deployed?
 - NAT
 - Network monitoring

Coming soon

- Internet design principles
- Class project overview