

Lecture #3: IPv6

WPI CS4516

Spring 2019

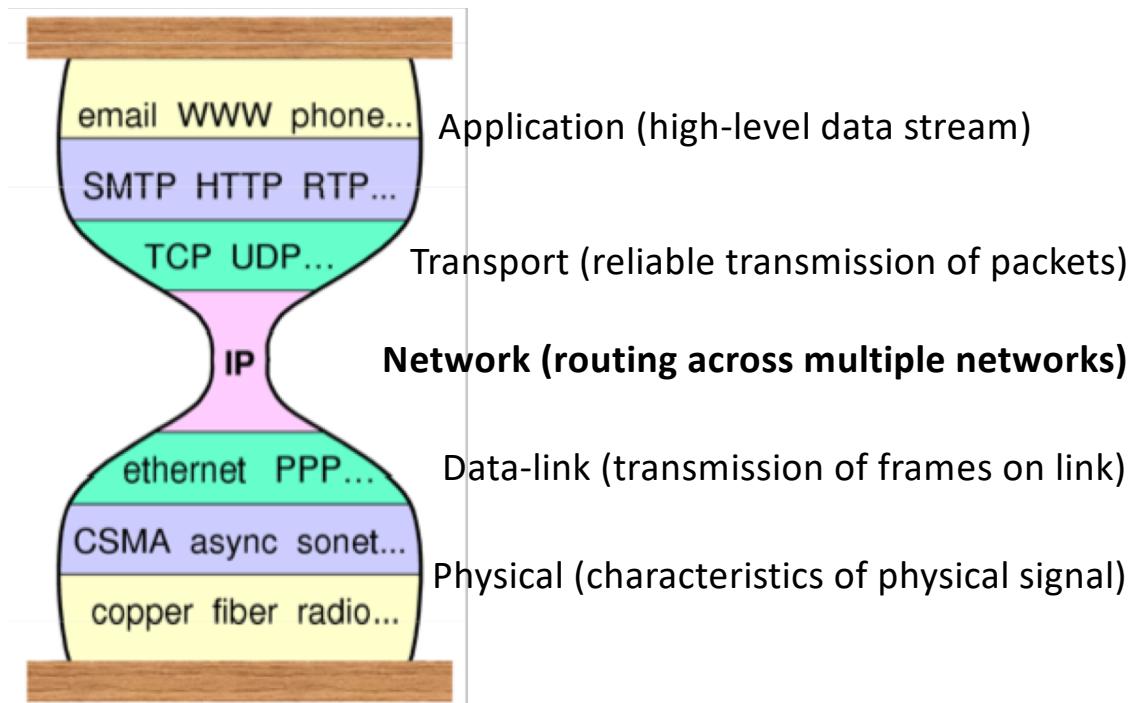
D term

Instructor: Lorenzo De Carli (ldecarli@wpi.edu)

*(slides include material from Christos Papadopoulos, CSU and
Craig Shue, WPI)*

IPv6

- **Network-level protocol** designed to be the successor of IPv4
- Responsibility: **datagram routing** across the Internet
- Implemented on **nodes and routers**



Background

- IPng recommended by IETF in July 1994
- Approved by Internet Engineering Steering Group in Nov. 1994
- Published as RFCs in December 1995.
- Formally assigned the name IPv6
- Article published June 1996

Motivation

- **Overcoming pressing IPv4 limitations** due to outdated design
- Introducing support for **new network-level primitives**
- What happened to **IPv5**?
 - Connection oriented experiment
 - Parts of it ended up in MPLS

A bit of history

IPv4

- Proposed in the early ‘80s
 - RFC 791, Sep 1981
- Deployed in 1983
 - Internet Flag Day, remember?
- **Still routes most internet traffic today**

It works... why change it?

- Multiple reasons
- The most pressing (and probably the most valid) one: **address exhaustion**
- IPv4 was designed in the early days of the Internet...
 - ...therefore 2^{32} addresses seemed more than enough!
 - But already in the early '90s it became clear that this was not the case

IPv4 address space limitation

- In the initial IP design, **every valid address** (i.e., minus reserved ones) is expected to be **globally routable**
 - Every device on the network must have a public IP
- Initially, addresses were partitioned as 8-bit network + 24-bit host address
 - **Extremely wasteful!**
 - 16M host per network

IPv4 class system (1984, RFC 791)

- **Class A:** 8+24, MSB is 0 (128 networks)
- **Class B:** 16+16, MSBs are 10 (16384 networks)
- **Class C:** 24+8, MSBs are 110 (2M networks)
- **Class D:** multicast, MSBs are 1110
- **Class E:** reserved, MSBs are 1111
 - Unusable as some network equipment will refuse to route it
- Designed to **give some flexibility in address assignment while keeping HW simple**

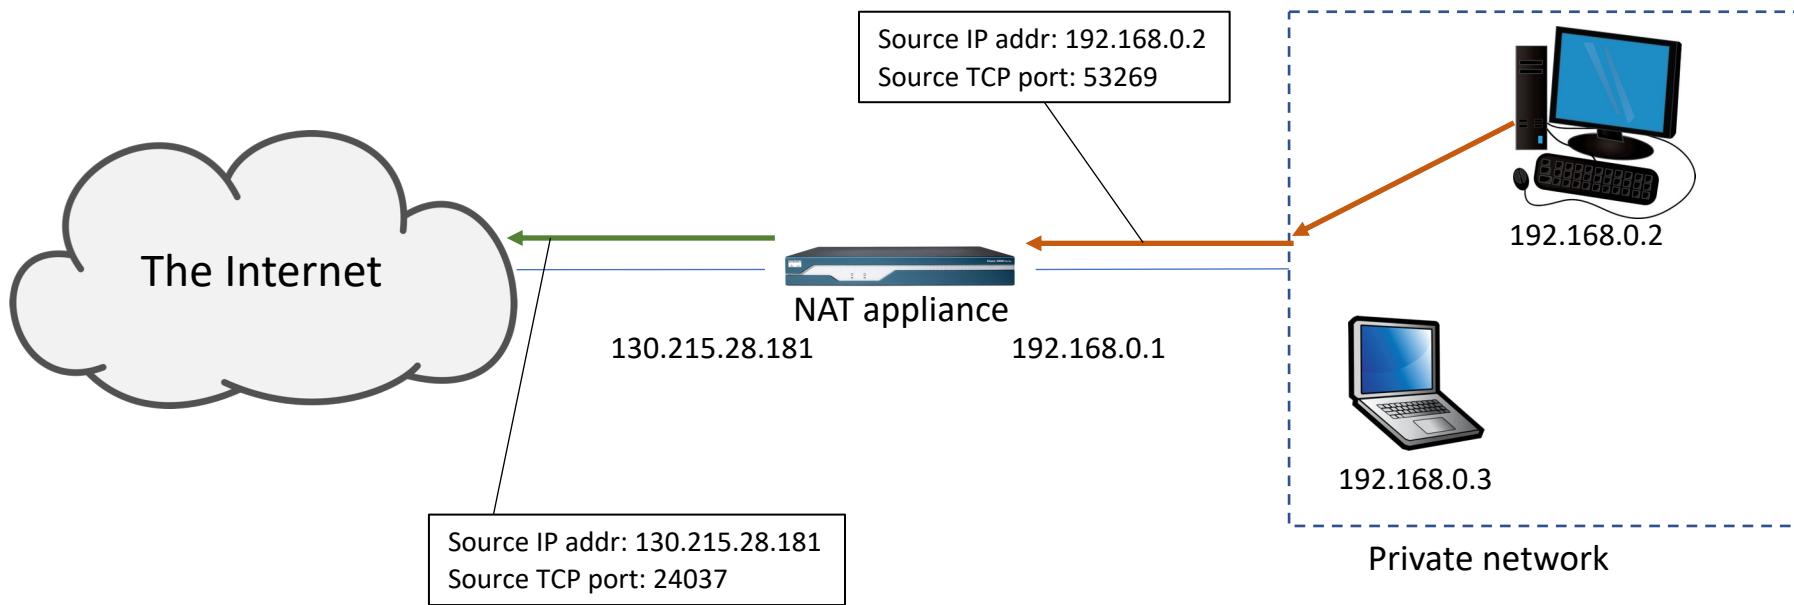
Classless IPv4

- By 1993, it was clear that **even the multi-class system was too wasteful**
 - The address space was used too sparsely and going towards exhaustion
- Solution: **classless IP (CIDR)**
 - Introduced by IETF in 1993
 - Does away with the notion of classes; **allows prefixes of any length to act as the subnet id**
 - Example: 130.215.28.0/24

NAT

- Classless IP works in tandem with the other “emergency measure” introduced to deal with address exhaustion: **NAT (RFC 1631, 1993)**
 - “Network address translation”
 - In practice, **breaks the Internet into a public scope and many private scopes**
 - Nodes in the public Internet must have public IP addresses
 - All private networks reuse the same blocks of private addresses
 - A dedicated network appliance (NAT) translates private addresses to public addresses (and vice-versa) when nodes in a private network communicate with the public Internet

NAT - II



Note: other types of NAT are possible

Types of NAT (RFC 3489, 3/2003)

- **Full-cone NAT:** 1-to-1 mapping between $(address, port_I)$ on internal node and $(address_G, port_G)$ on NAT gateway
 - Any external host can send packets to $(address, port_I)$ by addressing them to $(address_G, port_G)$
 - **Advantage:** hosts are reachable from the public Internet
 - **Disadvantage:** hosts are reachable from the public Internet

Types of NAT - II

- **Restricted-cone NAT:** 1-to-1 mapping between $(address_i, port_i)$ on internal node and $(address_G, port_G)$ on NAT gateway
 - An external host with IP $address_E$ can only send packets to an internal host only if the latter has previously sent a packet to $address_E$
 - “Hides” internal hosts behind the NAT

Types of NAT - III

- **Port-restricted cone NAT:** 1-to-1 mapping between $(address_I, port_I)$ on internal node and $(address_G, port_G)$ on NAT gateway
 - An external host with IP, port $(address_E, port_E)$ can only send packets to an internal host only if the latter has previously sent a packet to $(address_E, port_E)$

Types of NAT - IV

- **Symmetric NAT:** NAT gateway maps each flow (*address_S, port_S, address_E, port_E*) to a specific (*address_G, port_G*) on the gateway
- Note the difference between *cone NAT and symmetric NAT:
 - The former assigns **gateway address/port based on address source/port**
 - The latter assigns **gateway address/port by flow**
 - **Symmetric NAT** is the trickiest to bypass because it makes mapping unpredictable

NAT pros & cons

- When proposed, **NAT was controversial**
 - Breaks the principle “one node, one (or more) public IP”
- But it has advantages
 - **Decouples** the number of Internet-connected nodes in an organization from the size of its Internet presence
 - An organization (e.g., university) may want all computing nodes to have connectivity, but may only need a few IPs for public Internet presence, email etc.
 - **NAT can hide private network topology and make private node unreachable by default**

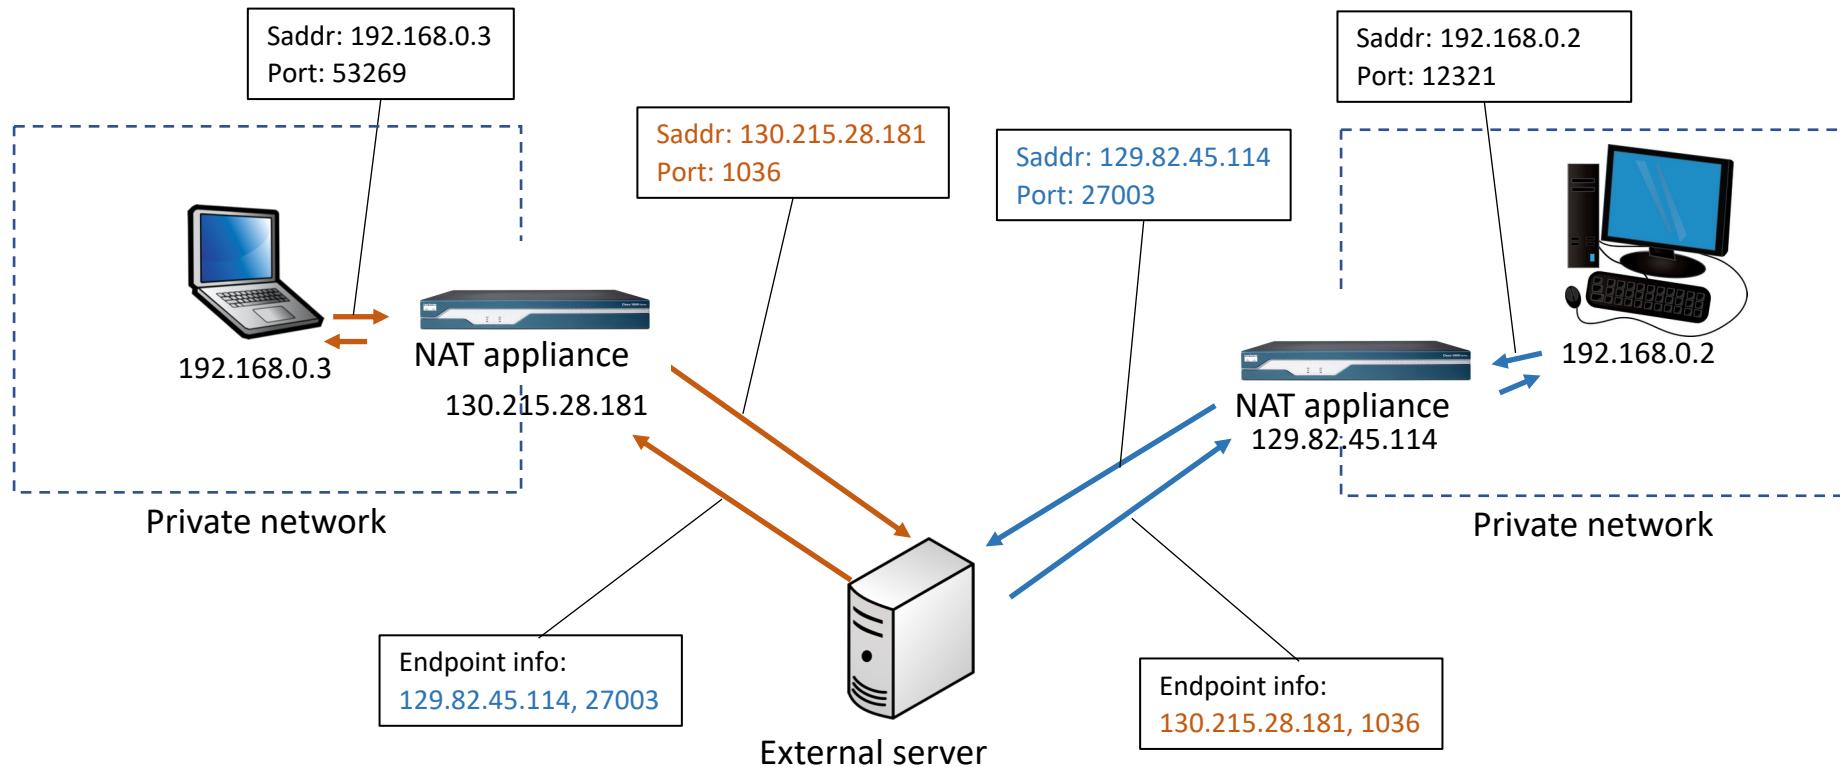
NAT issues

- Oftentimes, it is useful to “hide” a network behind a NAT but it is desirable to **allow some internal host addresses/ports to be reachable from the outside**
- **Various techniques:** hole-punching, port forwarding

NAT hole-punching

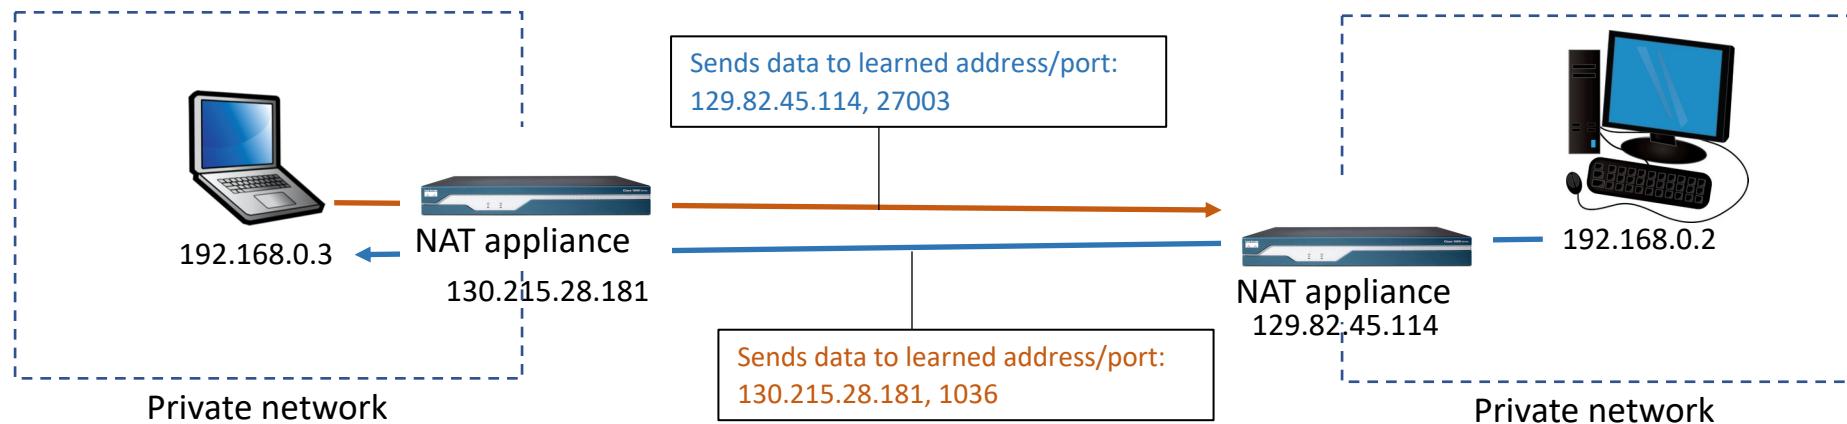
- Works for ***cone NAT**, where the association between internal and external addresses/ports is **persistent**
- Designed for situations where **two hosts behind different NATs want to communicate**
- Hosts use a third-party - a server w/ IP address – to determine the **external addresses and port to which are mapped**
- A form of this is defined in RFCs 3489/5389

Hole punching example



What's missing?

Hole punching example - II



If restricted NAT, will need to make multiple attempts (why?)

What if it does not work?

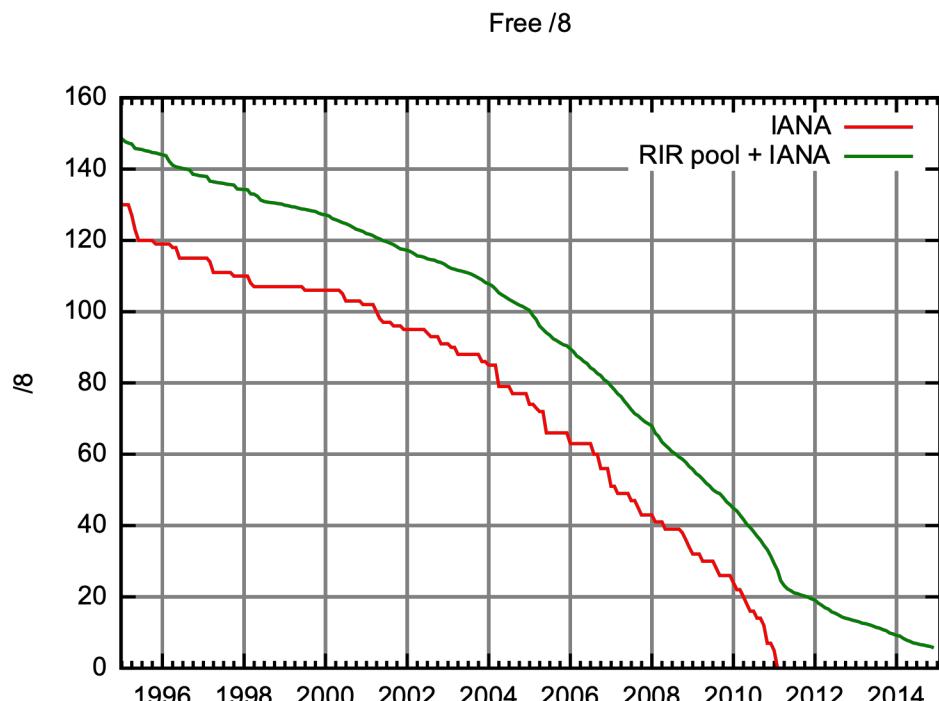
- (E.g.: symmetric NAT, nonstandard configurations, etc.)
- **Port forwarding:** configure an external port on the NAT gateway to **always forward traffic** to an internal host/port
- Those of you who are running their own home server probably do this ☺
- **uPNP** (supported by pretty much all home routers) automates port forwarding

CIDR + NAT

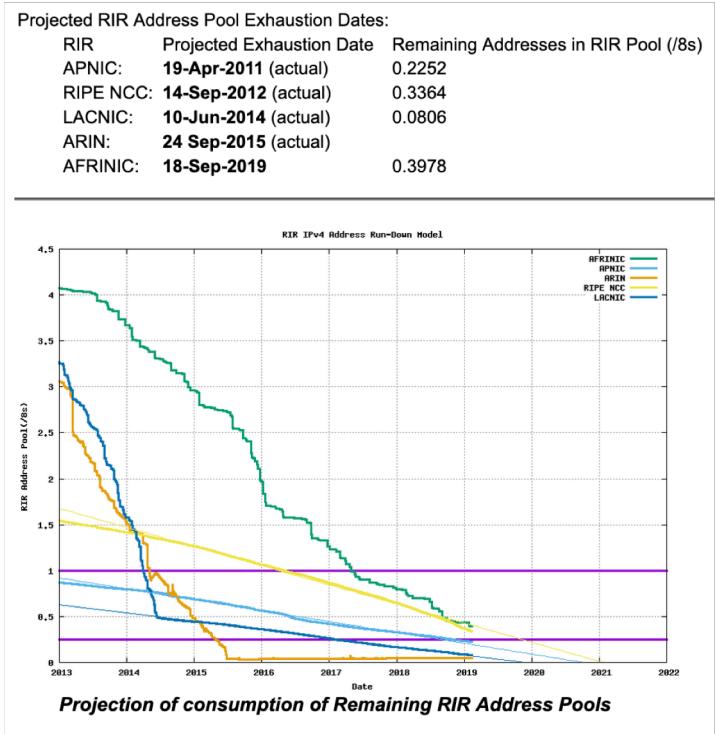
- Summarizing:
 - NAT **lowers the number of public IPs** that an organization needs to have Internet connectivity
 - CIDR **allows assignments of very small blocks of public IPs** to an organization
 - By allowing many nodes to use private addresses, **the pressure on the IP address pool was drastically decreased!**

Yes, but did it work?

- Quite well, in fact – but we are at the end



By Mro - Own work, CC BY-SA 4.0,
<https://commons.wikimedia.org/w/index.php?curid=10593349>



Now, let's move on to IPv6

IPv6 design principles

- **More addresses!**
 - Motivation is less obvious than it may seem – we'll go over it
 - **Also, more expressive** addressing scheme
 - IPv6 implements several primitives that in IPv4 require external protocols
- **Ease of transition** from IPv4
- **Better security**
- Better QoS, support for host mobility, etc.

Why more addresses?

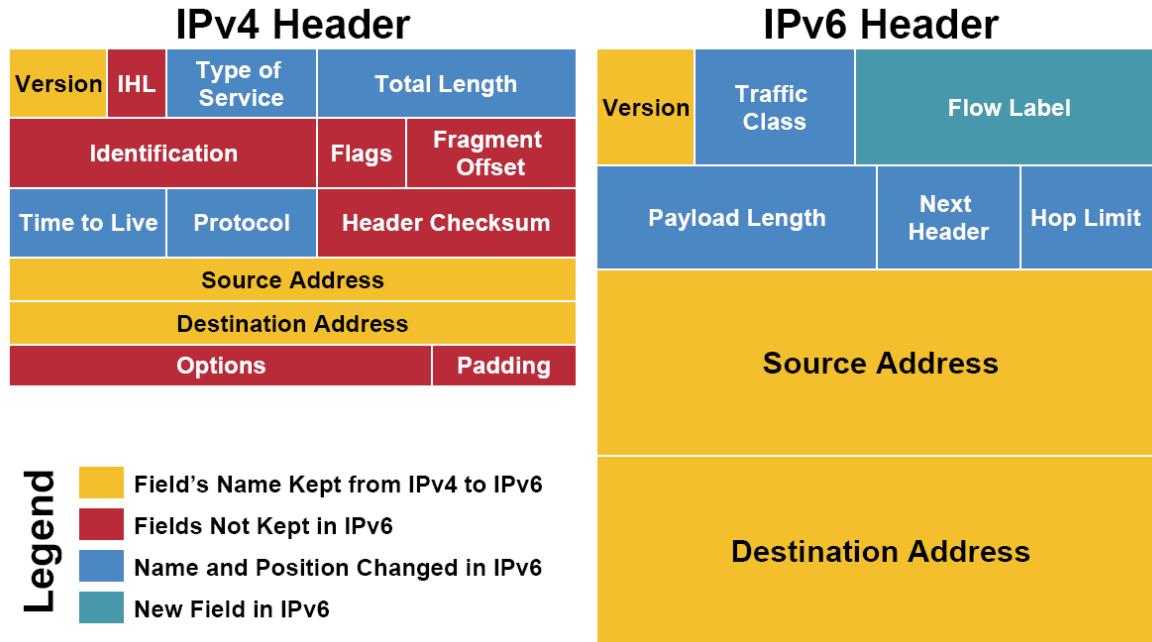
- One reason is obvious: **IPv4 address space exhaustion**
- But there is also a vision on the **future of computing** (1996)
 - General-purpose computer market will keep expanding
 - “Nomadic personal computing devices”
 - “Networked entertainment”
 - “Device control”



IPv6 Header



Compare to IPv4 header

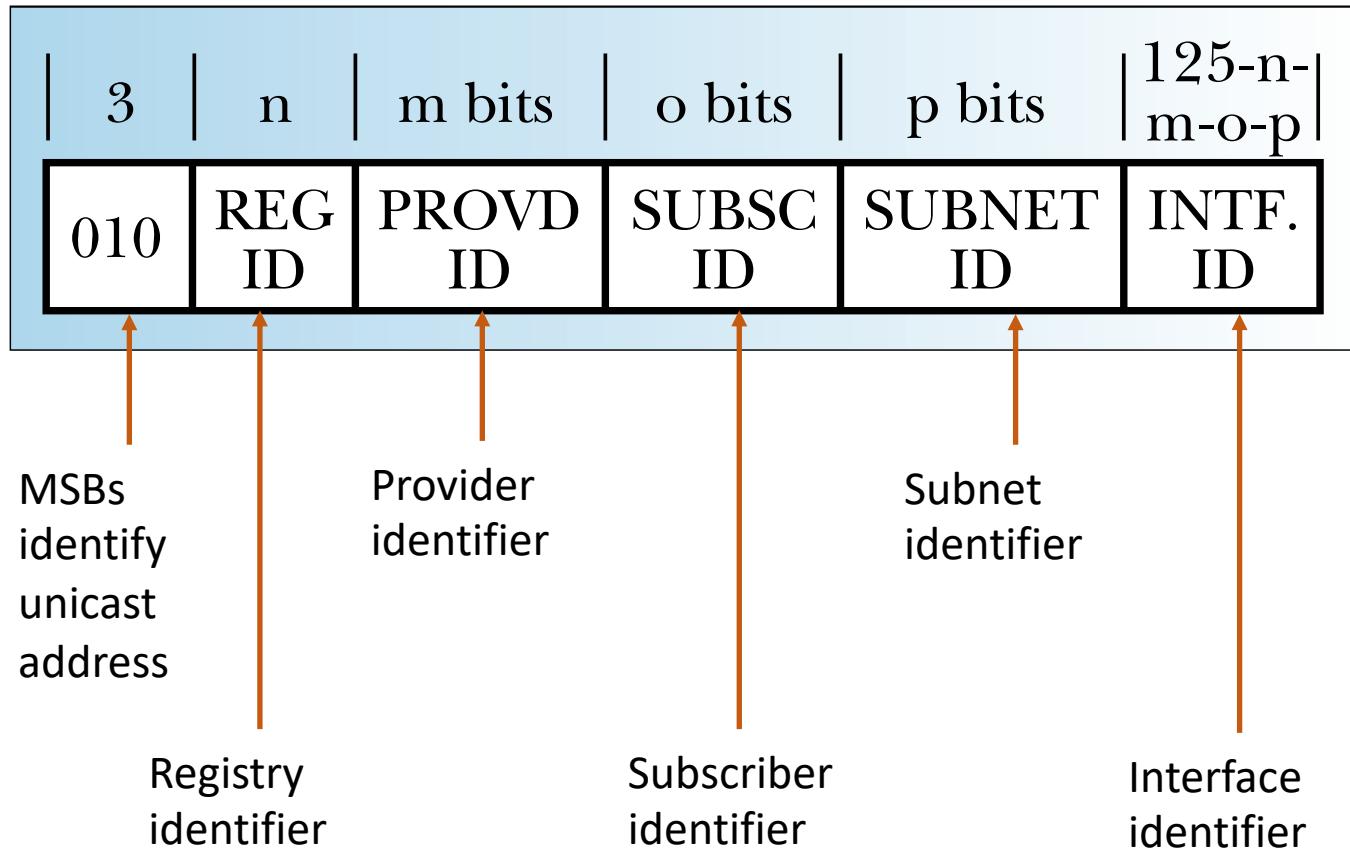


- Addresses are 4 times longer (3.4E38 possible values of an address field)
- However, header is only 2 times longer (attempt to keep overhead acceptable)
- (Still, it made the life of router designers quite difficult)

IPv6 Address Space Allocation

- Vast majority is **reserved** for future use
- **Aggregatable Global Unicast Addresses** make up **1/8** of all addresses
- “Link local use” addresses are **usable on local network**, but **not globally unique**
 - Autoconfiguration
- “**Site local use**” is used by **private networks**

Global unicast addresses



Local-use addresses

- Designed to be in large part **autogenerated** without the need for external information
 - Enable a node to get an address guaranteed to be (at least locally) unique without relying on other nodes
 - Useful for node **autoconfiguration** when first connected to the network
- How does it work?

Local-use addresses - II

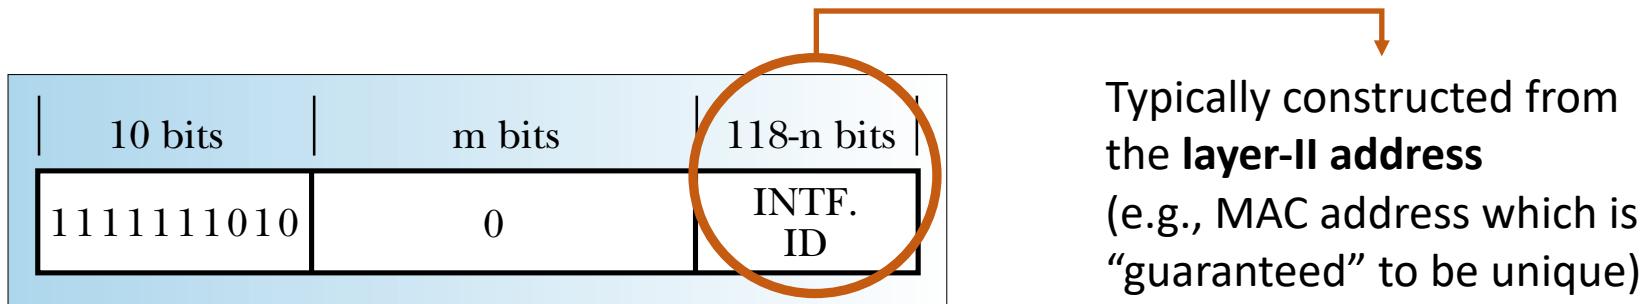


Figure 4. Link-Local address format

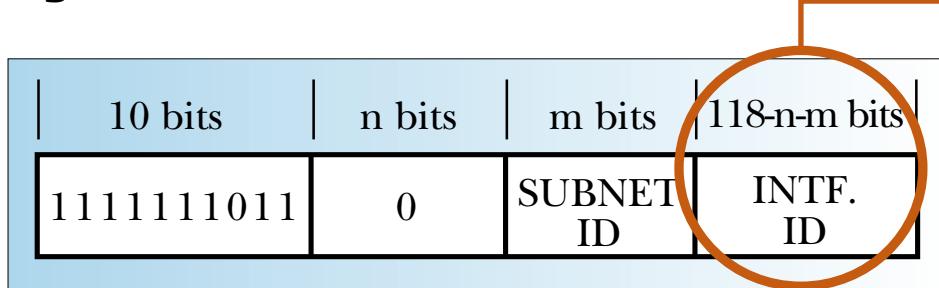


Figure 5. Site-Local address format

Goal: ease of transition

- It was obvious to IPv6 designers that the Internet could not be coerced to switch from IPv4 to IPv6
- In the early 90's, the network was already large and complex enough that **a flag-day style transition would have been untenable**
 - Who remembers the other flag day?
- **Significant capital investment** in IPv4 hardware, software stack, etc.
- IPv6 is designed from the ground up to **coexist** with IPv4 and **make the transition as painless as possible**

IPv4-compatible addressing scheme

1: IPv4-mapped IPv6 addresses

- First, IPv6 address can **embed** IPv4 addresses

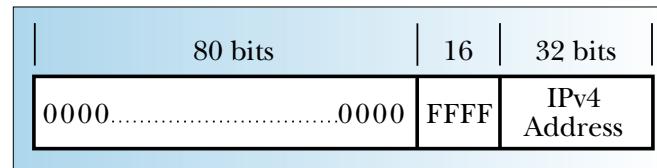
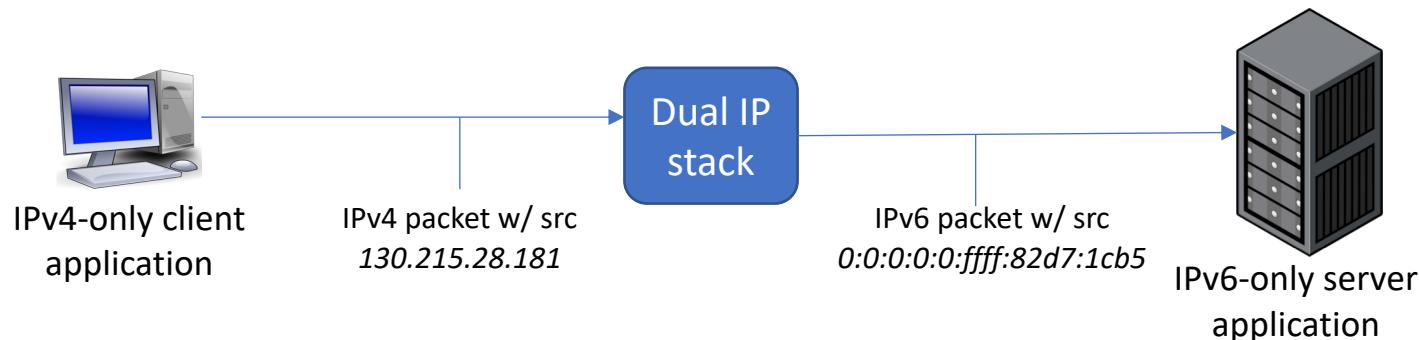


Figure 7. IPv4-mapped IPv6 address

- What is this for? (RFC 4038, “Application Aspects of IPv6 Transition”)



2. IPv4-compatible IPv6 addresses

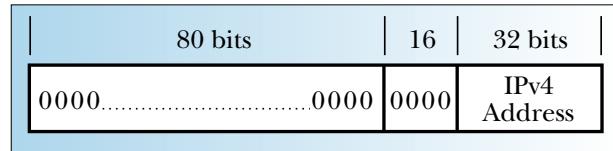
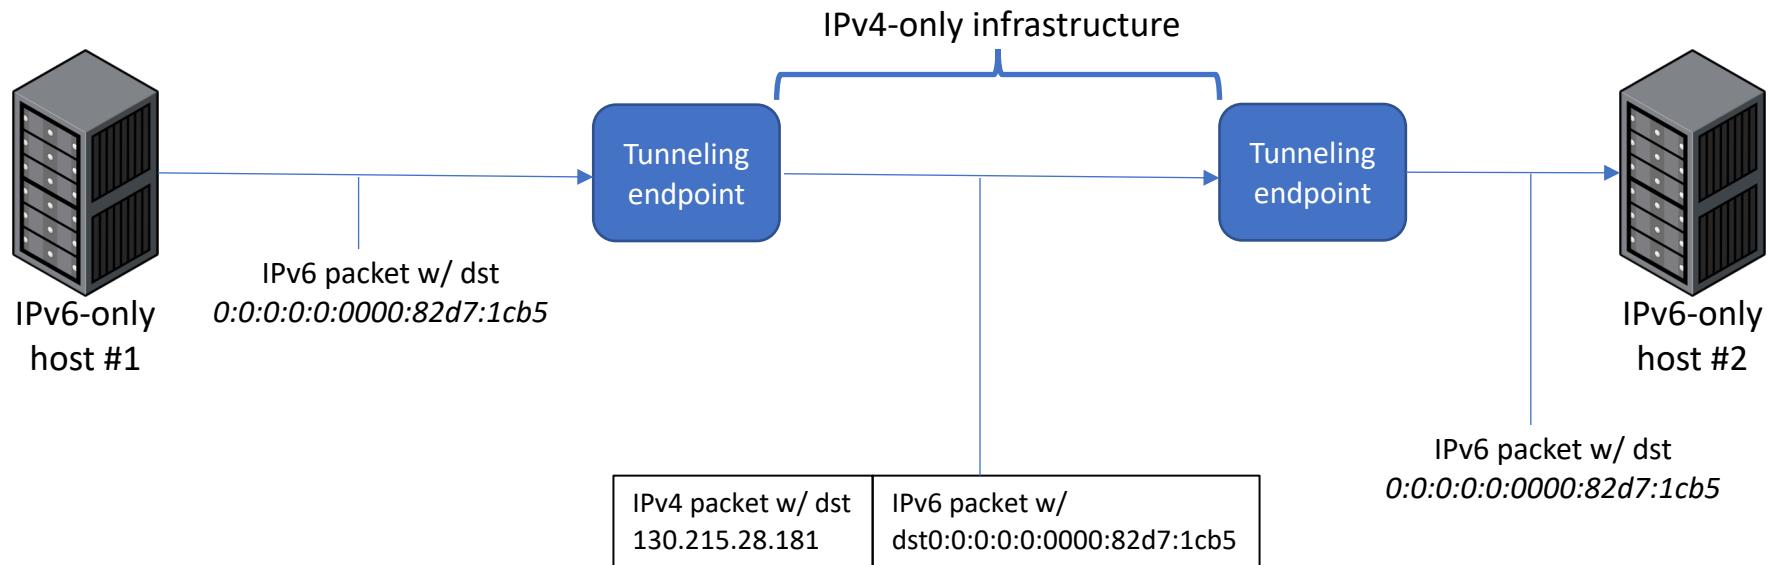


Figure 6. IPv4-compatible IPv6 address

- These ones solves an **infrastructure issue**
 - How to route packets between two IPv6 hosts **through IPv4 infrastructure**
 - Idea:
 - **Tunnel IPv6 in IPv4**
 - Derive destination IPv4 address from destination IPv6 address

IPv6 over IPv4 tunneling

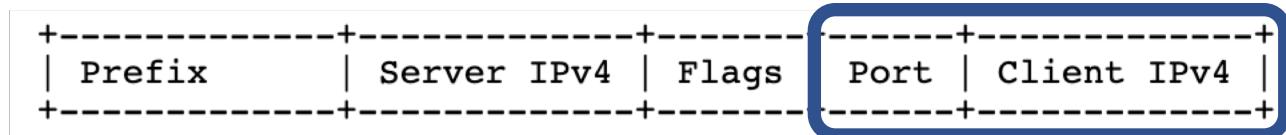


Teredo tunneling

- Another flavor of IPv6-over-IPv4 tunneling
- IPv6 tunneled in UDP over IPv4
- Why? Regular tunneling hides layer 4 and therefore breaks NAT hole punching & related techniques
- Developed at Microsoft, standardized in RFC 4380

How does Teredo work?

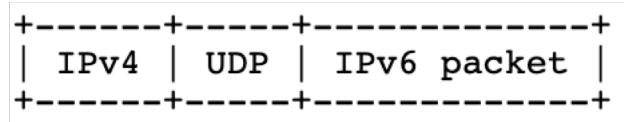
- Client interacts with a **Teredo server** to find out **mapping** between local and NAT address/port (a form of hole punching)
- Server generates **an IPv6 address for the client**. The address encodes the public address/port used on the NAT gateway:



Public IPv4 address/port for
the IPv6 client behind NAT

How does Teredo work? - II

- After client is assigned an IPv6 address, it uses a **Teredo relay** to communicate with other IPv6 peers
- The client **encapsulates** IPv6 packets into UDP over IPv4:



- The relay receives these packets, **decapsulates them**, and sends them to the IPv6 destination
- When receiving packets for the host, the relay performs **encapsulation**

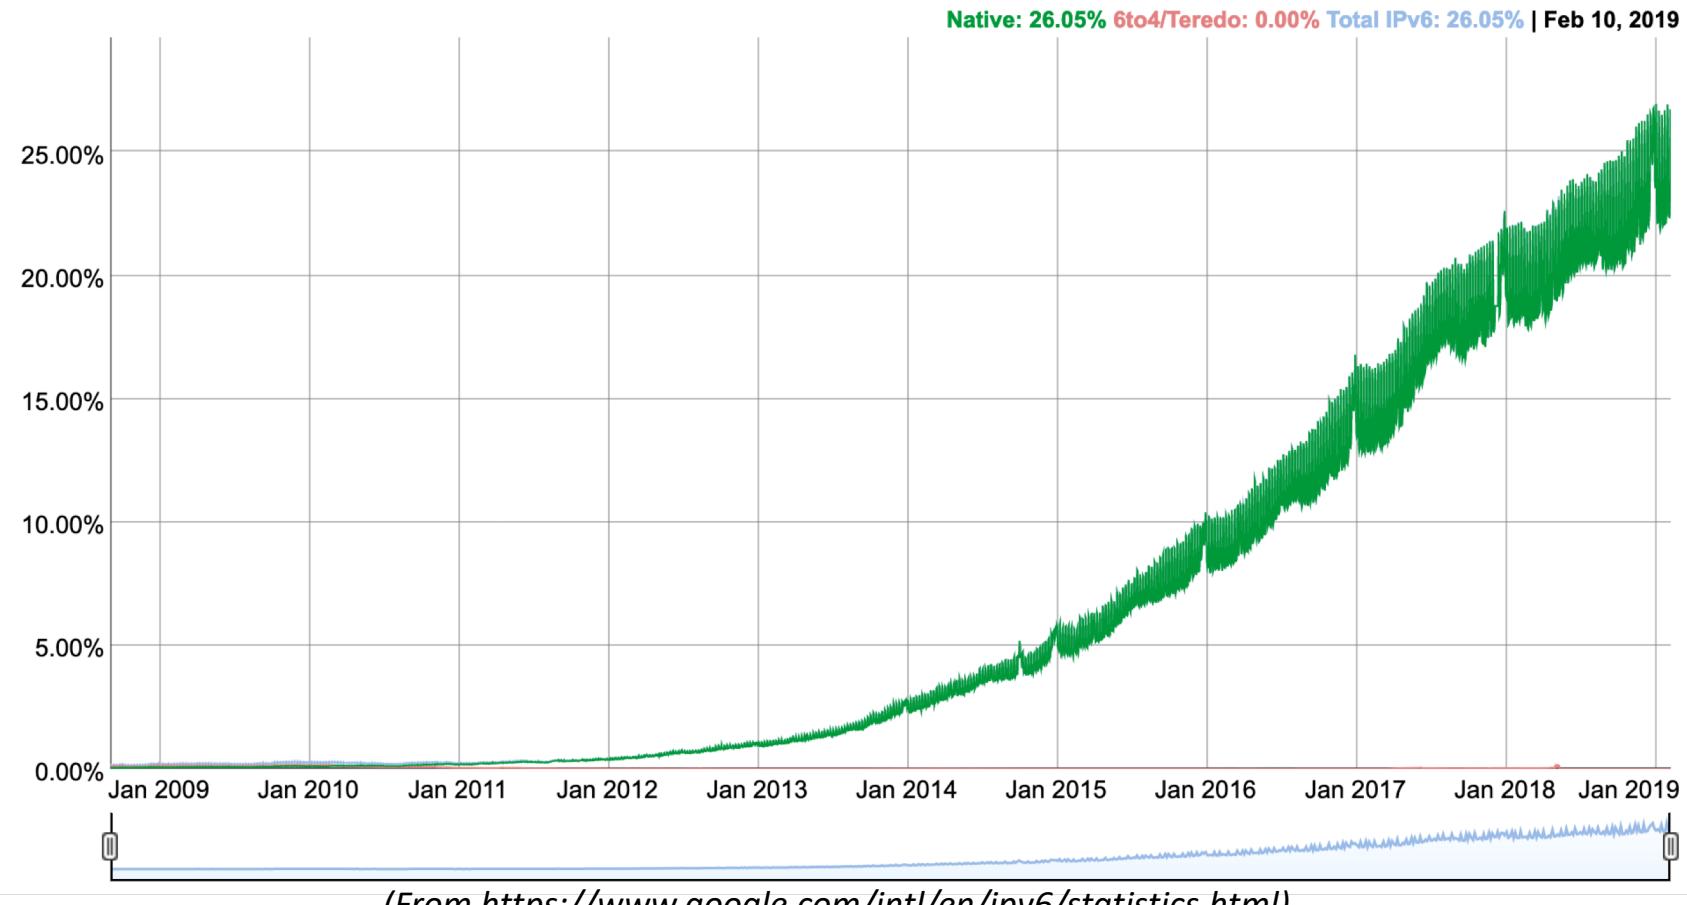
Other tunneling type

- Over time, various **other tunneling techniques** have been defined/used
 - 6rd
 - ISATAP
 - ...
- All these technologies are variations of the tunneling principle with modifications to solve specific problems
- Bottom line: **the IPv6 community went to great lengths to ensure smooth transition**

Did they succeed?

IPv6 Adoption

We are continuously measuring the availability of IPv6 connectivity among Google users. The graph shows the percentage of users that access Google over IPv6.



Encouraging IPv6 adoption

- World IPv6 day: 6/8/2011
 - Several large companies (Facebook, Google, Akamai, ...) enabled IPv6 on their public-facing network for 24 hours
- World IPv6 launch: 6/6/2012
 - Make the activation permanent

Adoption problems

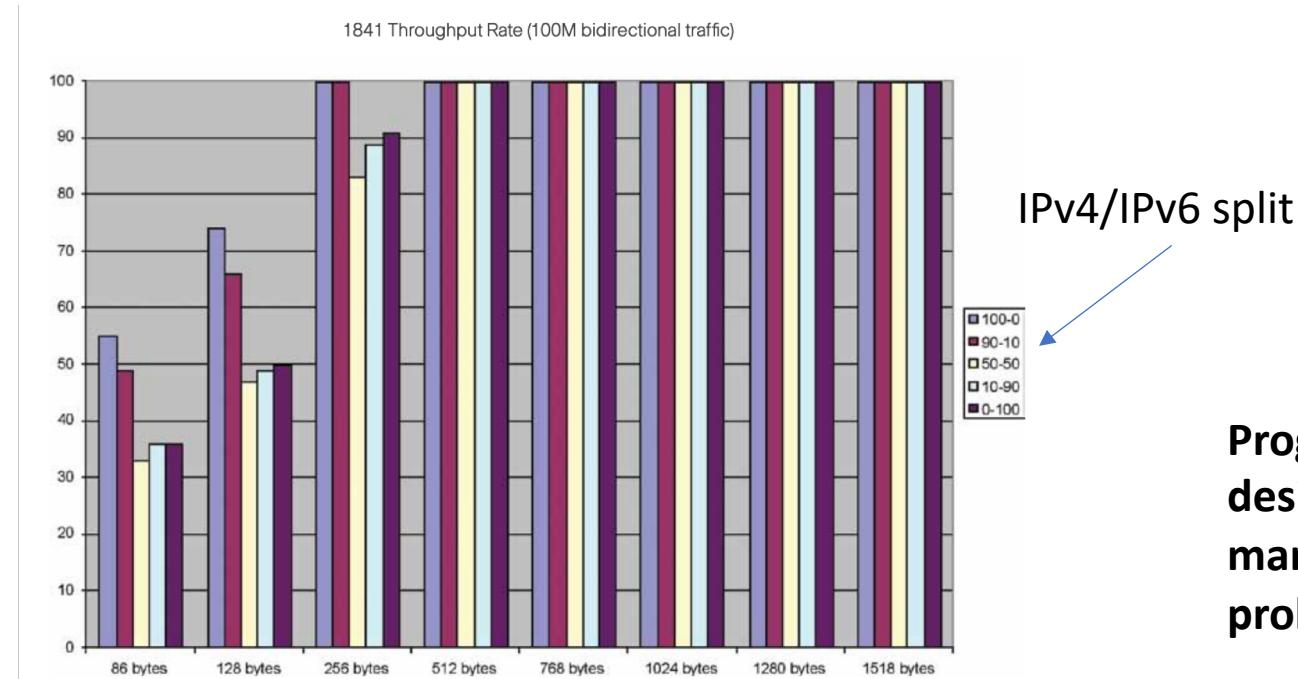
- **Malfunctions**
 - A typical OS networking stack is designed to prefer IPv6 connectivity when available
 - Especially in the early day of IPv6 adoption, a network may advertise IPv6 connectivity but **have it configured incorrectly**
 - **Consequence:** a node may attempt to sends packets via a **broken IPv6 configuration**, rather than a **working IPv4 one**
 - Rare today due to **increased operational experience**/familiarity of operators with IPv6

Adoption problems - II

- Any respectable router implements IP forwarding using **custom hardware logic**
- IPv4 addresses are **substantially different** (shorter, w/ different internal structure) than IPv6 ones
- Implementing IPv6 forwarding requires **redesigned hardware**
 - Early dual IPv4/IPv6 routers implemented IPv4 forwarding in hardware, IPv6 forwarding in software, with **substantial performance drop**

Adoption problems - III

- Even when a router is designed from the ground up to support IPv6, it may be difficult to match performance of IPv4 forwarding
 - IPv6 address lookup requires larger buffers, more operations
 - E.g.: 2007 comparison for CISCO 1841 router:



Goal: better security

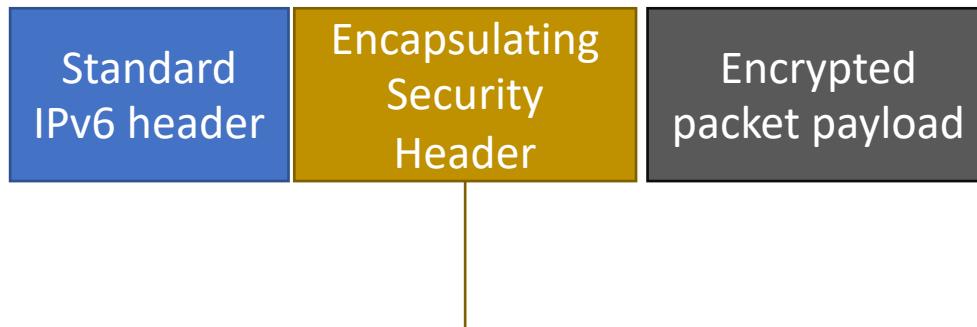
- IPv4 does not have any **security mechanism** built-in
 - Security properties built on top if it
- IPv6 has **IPSec-style security** built in
- What is IPSec?
 - Protocol providing:
 - **Authentication** (i.e. proof of origin) and **integrity** using the IP Authentication Header (AH)
 - **Confidentiality** (encryption) using Encapsulating Security Payload (ESP)

Authentication/Integrity in IPv6



- Covers all immutable fields of IPv6 header (e.g., exclude flow label)
- Includes:
 - Various metadata
 - A sequence number to prevent replay attacks
 - An integrity check vector (hash, keyed hash etc.)

Confidentiality in IPv6



- Provides metadata describing specifics of encryption used on the payload
- Includes:
 - A session identifier (security association identifier)
 - Sequence number

But wait... there is more

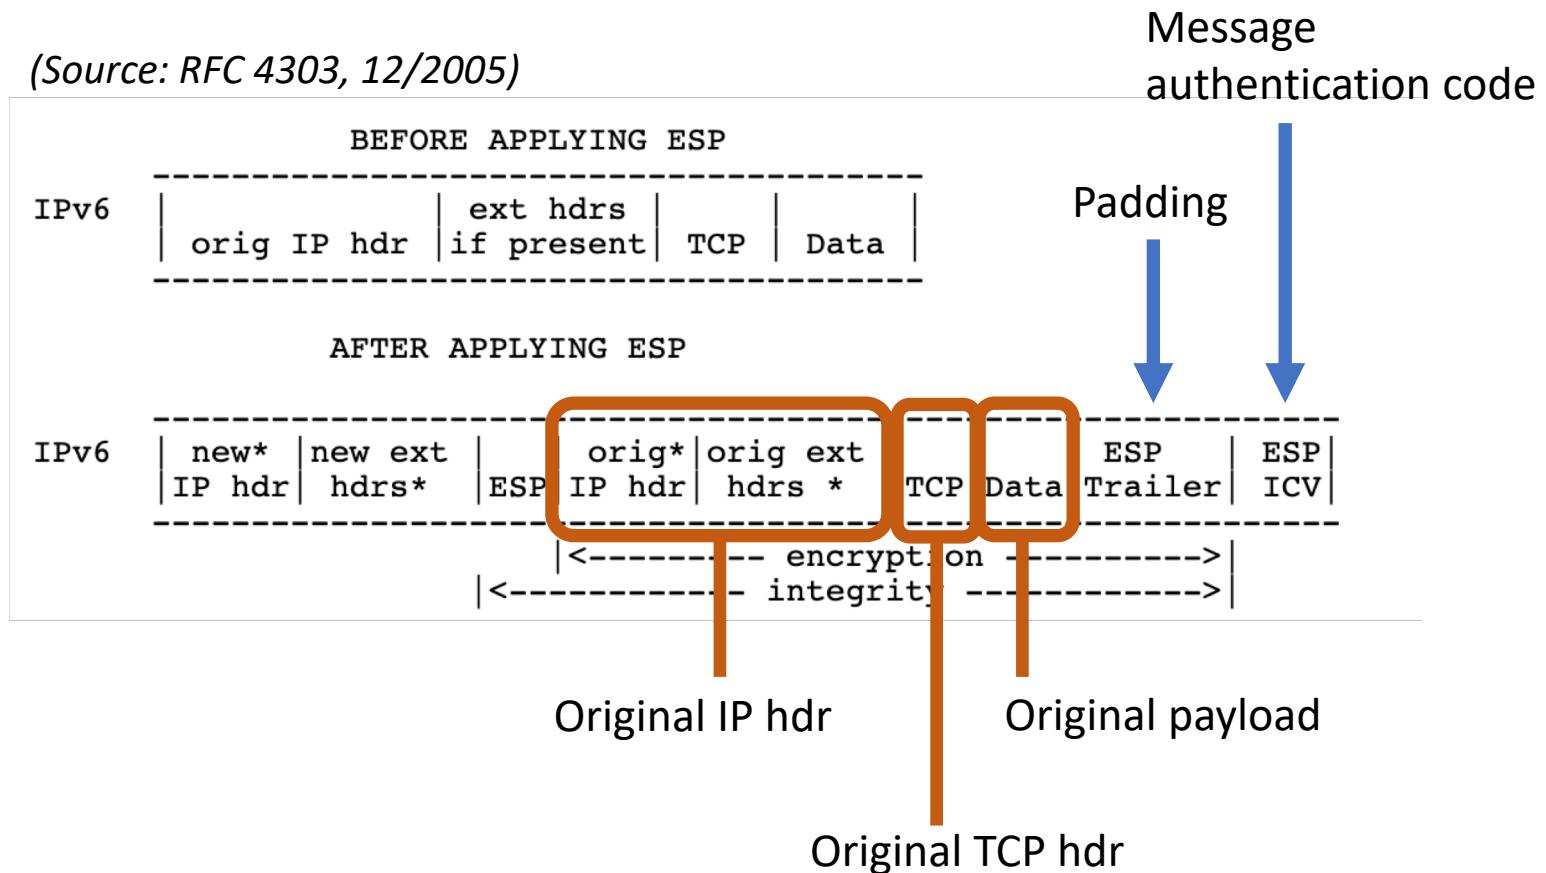
- What we have seen refers to authentication/confidentiality in **transport mode**
- IPv6 also support secure communication in **tunnel mode** (AKA VPN)

IPv6 tunnel mode

- **Tunnel mode** allows to tunnel IPv6 packets within IPv6 packets. Security can be applied so that the ESP header covers the entirety of the tunneled packets (including addresses and other header fields)
 - **Advantage:** source and destination (not just content) remain secret!

IPv6 tunnel mode - II

(Source: RFC 4303, 12/2005)



Other IPv6 goals

- Fix a bunch of IPv4 shortcomings
 - Better QoS
 - Better multicast/anycast
 - Better option handling

IPv6 QoS/flow labeling

- Providing quality of service at network level it is always been a challenging task
 - **Integrated services** (intserv): per-flow QoS
 - **Differentiated services** (difserv): per-class QoS
- IPv6 provides **two mechanisms**:
 - **Flow labels**, allowing to mark and identify set of flows that should receive a certain treatment
 - Should be randomly assigned, so that they can act as indices in hash tables
 - Do not specify **how** a flow should be treated – that must be determined using a different protocol
 - **Priorities among traffic classes**

IPv6 multicast

- Propagating packets to **multiple destinations** has obvious benefits for **one-to-many communication**
- IPv4 includes:
 - **Broadcast**: transmit to all hosts on a subnet
 - **Multicast**: transmit to specific sets of hosts
 - Both use sets of reserved addresses
- IPv6 works in the same way but **drops broadcast addresses**
 - Why?
 - Limited usefulness, redundant with layer II, obvious security issues

IPv6 anycast

- Send packets **to one among multiple possible destinations**
- Use cases?
 - **Mobility**
 - **Load-balancing**
- No reserved addresses – any address from the unicast address space can be used for anycast

Better option handling

- **Options** are used in IPv4 and IPv6 to **provide functionality beyond basic routing**
- In IPv4
 - **Limits** on individual and total **option length**
 - Option support required
- In IPv6:
 - “Next header field” points either to layer-IV packet or option, specifying its type
 - **No limit on number of options**, higher limits on length
 - Routers can skip options they don’t recognize/care about