

# Lecture #10: Software-defined Networking

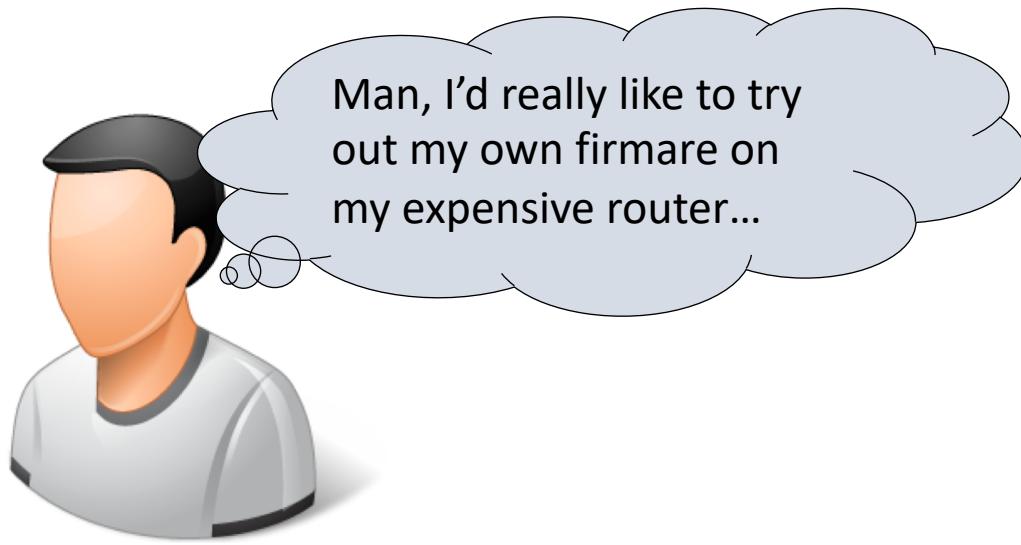
WPI CS4516      Spring 2019      D term

*Instructor: Lorenzo De Carli ([ldecarli@wpi.edu](mailto:ldecarli@wpi.edu))  
(slides include material from Christos Papadopoulos, CSU  
and Craig Shue, WPI)*

# What is this lecture about?

- OpenFlow: a paradigm to **simplify network management** by **centralizing control**
- We will review:
  - History & motivation
  - How OpenFlow works
  - Possible applications
  - Problems & possible solutions

# Motivation



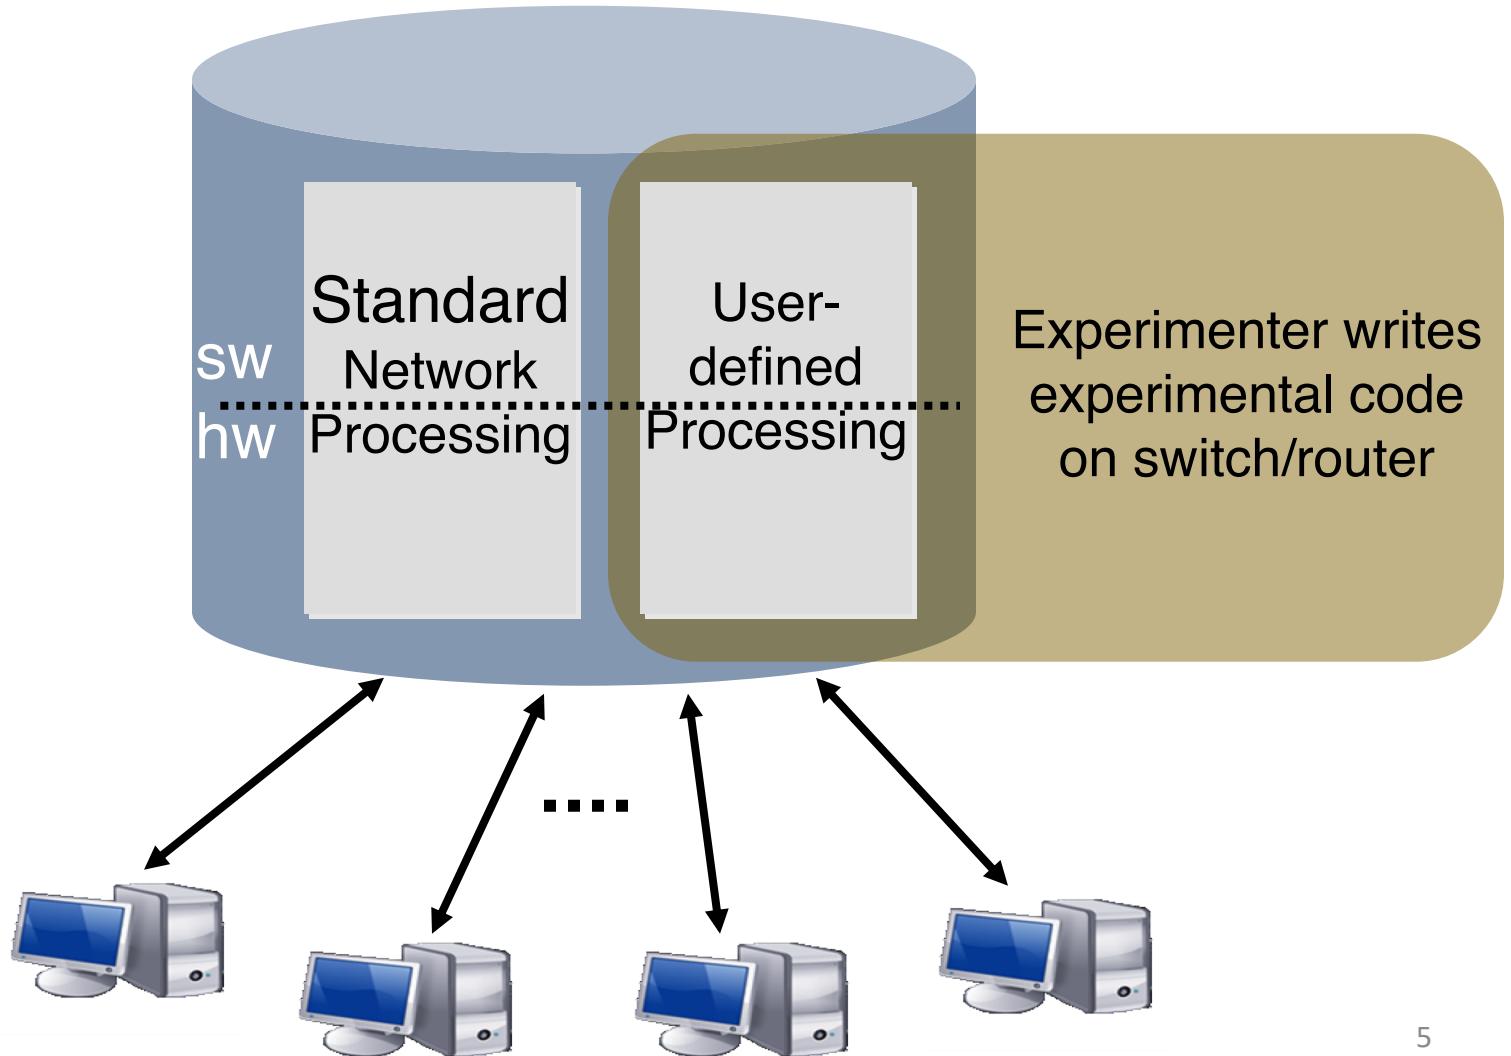
**Networking  
researcher**



# Motivation



# Experimenter's Dream



# Options?

- Commercial vendor **won't open software and hardware development environment**
  - Complexity of support
  - Market protection and barrier to entry
- Hard to build my own
  - **Prototypes are flakey**
  - Software only: **Too slow**
  - Hardware/software: **Fanout too small**  
(need >100 ports for wiring closet)

# Some research equipment pre-dating OpenFlow

- **Click (Kohler et al. '99): modular software router**
  - Lego approach to software router
  - Composition of modules offering various functions (packet routing, scheduling, classification, ...)
  - **Issue:** software routing on PC is slow
- **NetFPGA (Naous et al. '08): reconfigurable FGPA for networking**
  - Standard FPGA w/ Ethernet interfaces and a set of reusable packet processing modules
  - **Issue:** only four network interfaces (not enough for realistic switching/routing tasks!)

# More on the Wish List

- **Isolation:** Regular production traffic untouched
- **Virtualized and programmable:** Different flows processed in different ways
- Equipment we can **trust** in our wiring closet
- **Open development environment** for all researchers (e.g. Linux, Verilog, etc).
- **Flexible definitions of a flow**
  - Individual application traffic
  - Aggregated flows
  - Alternatives to IP running side-by-side
  - ...

# In summary...

- **Research** requires **flexible networking equipment**
- **Custom-built solutions** (software, hardware) are **too limited** for many applications
- **Adding flexibility to commercial switching/routing equipment** is a much more powerful solution
  - OpenFlow was created to achieve this!

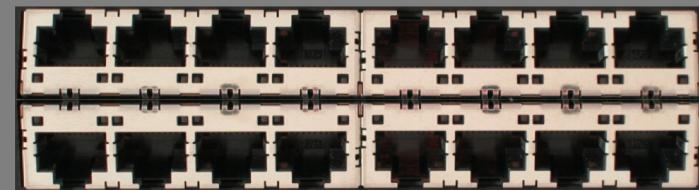
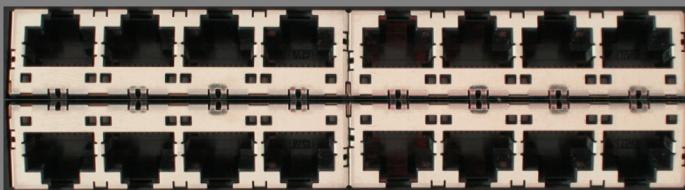
# Is this all there is to know?

- **No!**
- OpenFlow started as an academic project but found plenty of industrial applications!
- But let's begin from the beginning...

# How does OpenFlow work?

# Today's Switching Hardware

Ethernet  
Switch/Router



# Logical Division in Networking

Control Plane (Software)

---

Data Plane (Hardware)

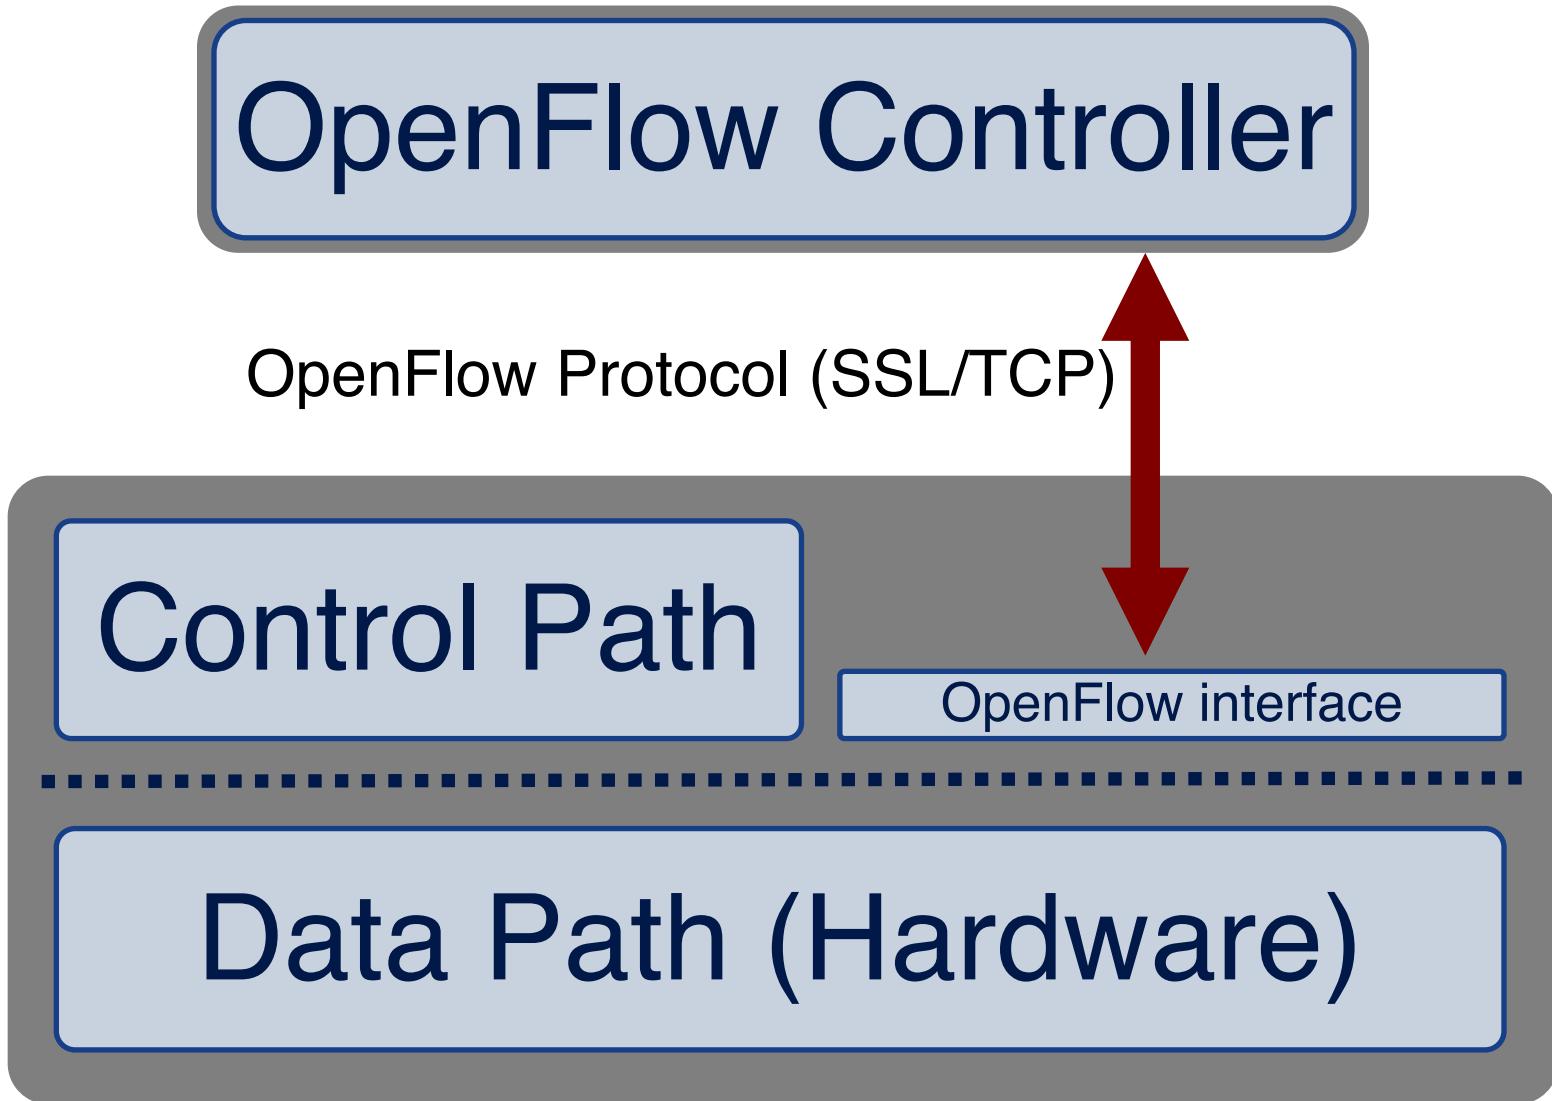
**Can someone explain what this means?**

# Traditional switch

Control Path

Data Path (Hardware)

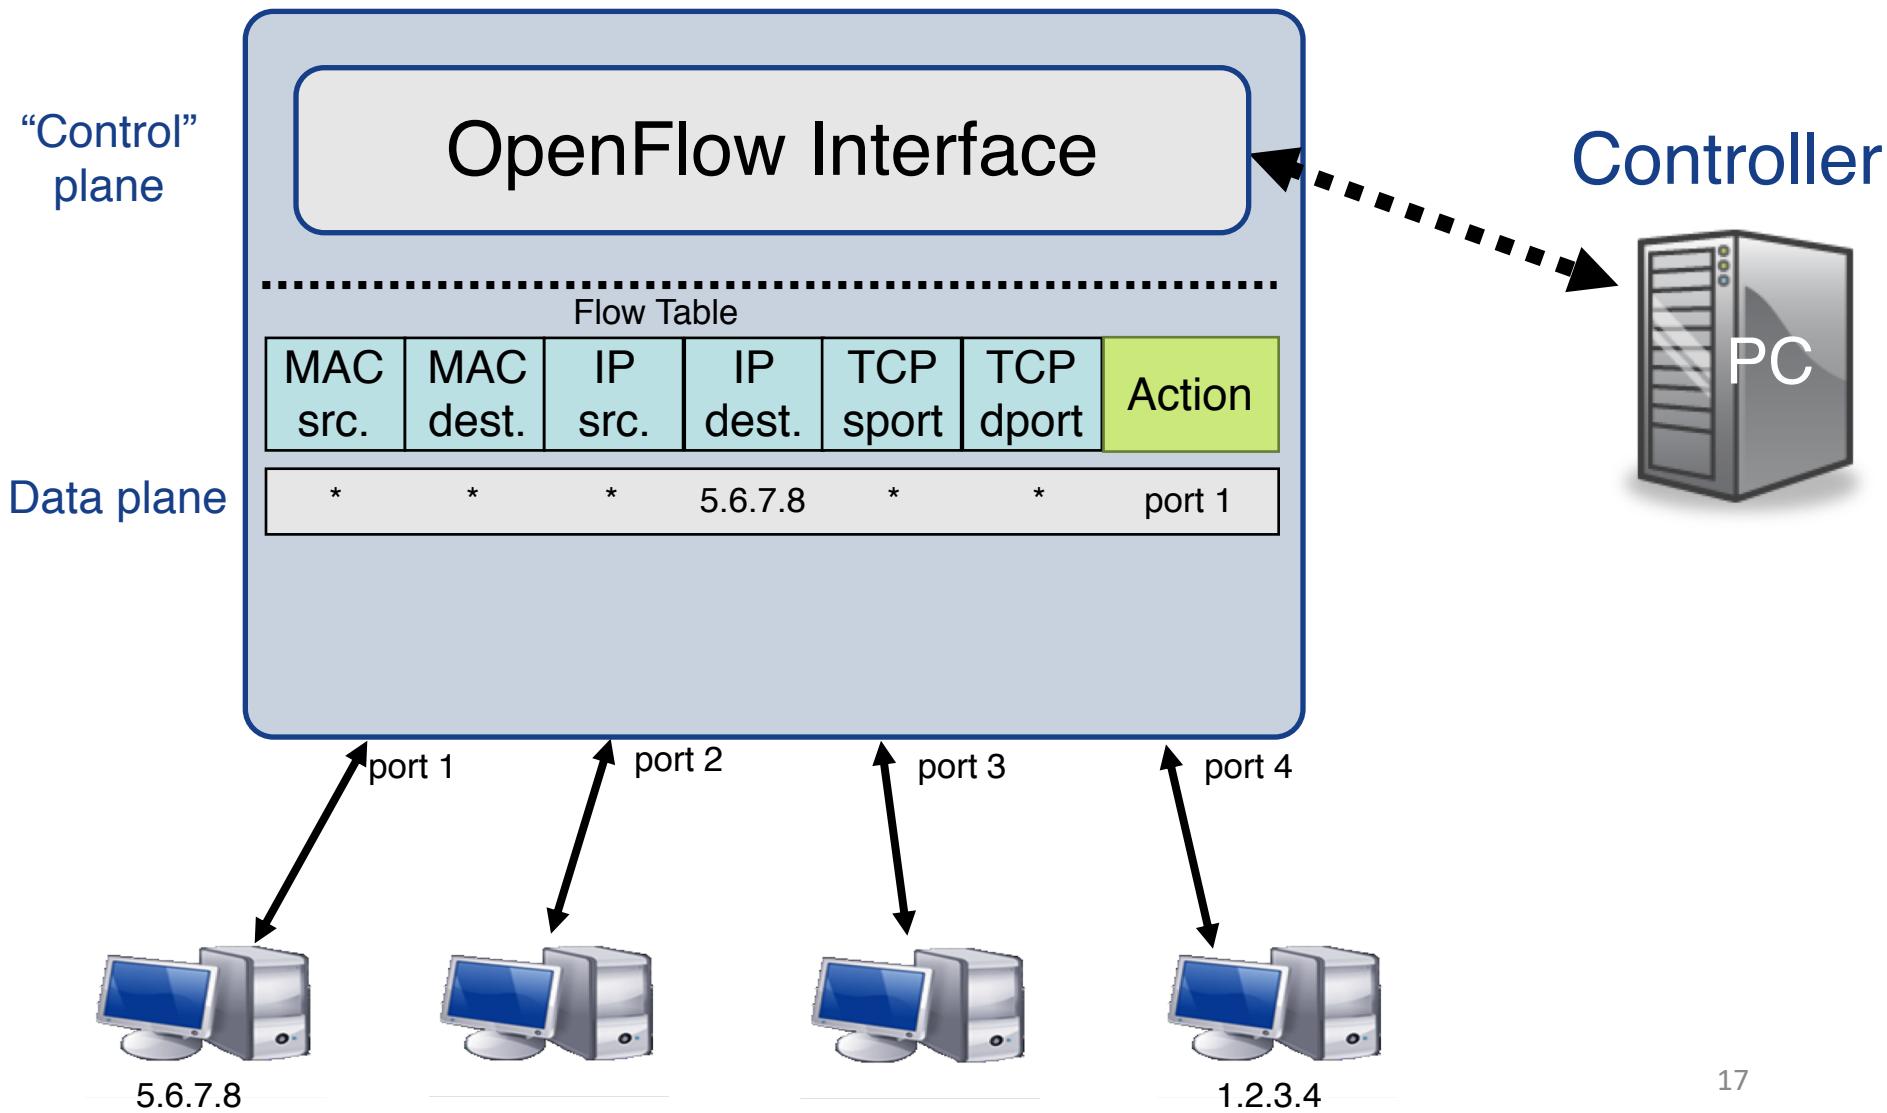
# Openflow switch



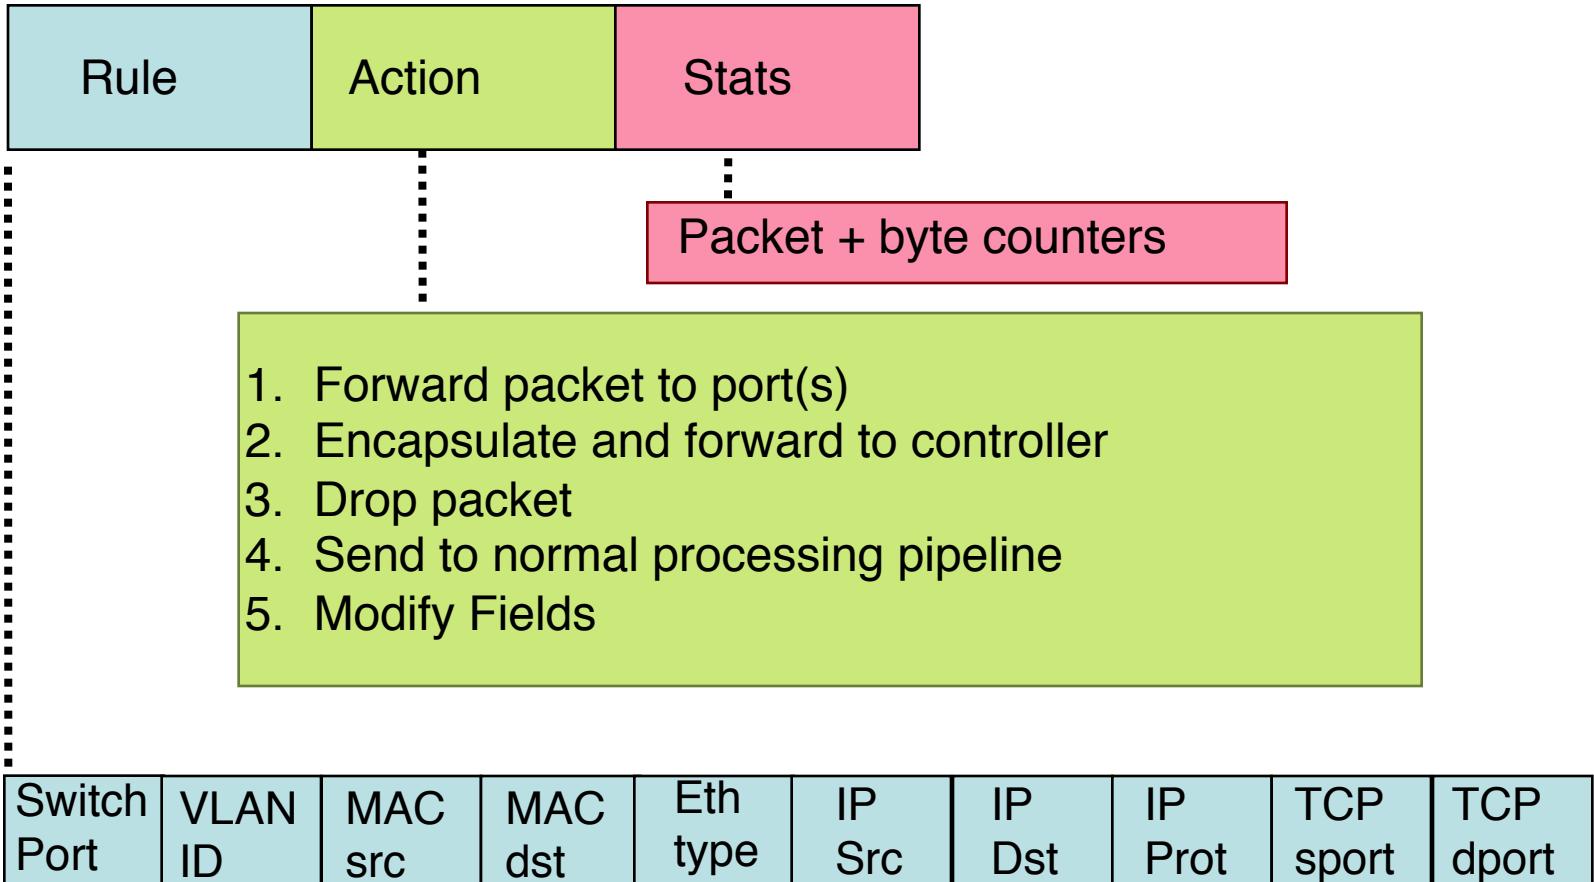
# OpenFlow principles

- A traditional switch implements:
  - The **data plane** logic to forward frames
  - The **control plane** logic to decide how to forward frames (forwarding tables, spanning tree, etc.)
- An OpenFlow switch implements:
  - A somewhat **flexible data plane logic** supporting various types of rules
  - **No control plane** – all decisions are delegated to an OpenFlow controller
  - Typically OpenFlow controller installs **rules** in switches, so that it does not have to analyze every packet

# OpenFlow model



# Flow Table Entries



+ mask what fields to match

# Examples

## Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f...	*	*	*	*	*	*	*	port6

## Flow Switching

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
port3	00:20..	00:1f..	0800	vlan1	1.2.3.4	5.6.7.8	4	17264	80	port6

## Firewall

Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Forward
*	*	*	*	*	*	*	*	*	22	drop

# Examples

## Routing

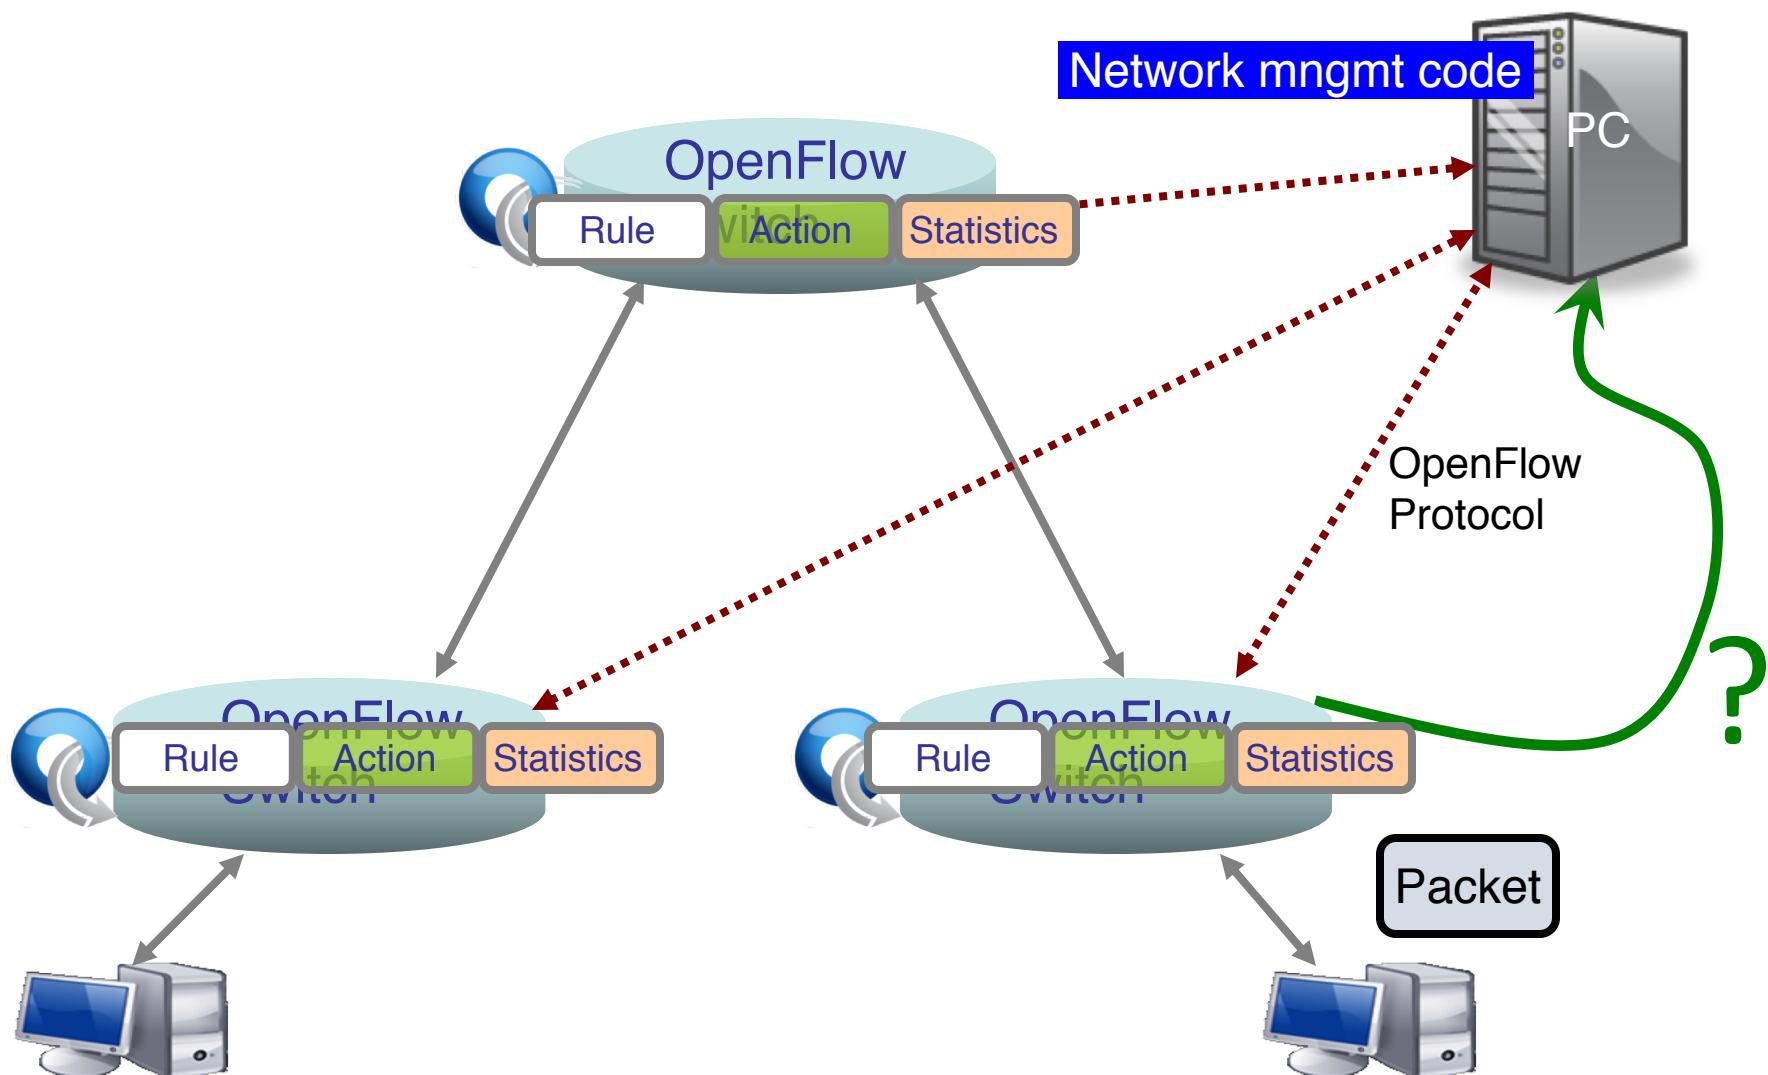
Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	*	*	*	*	5.6.7.8	*	*	*	port6

## VLAN Switching

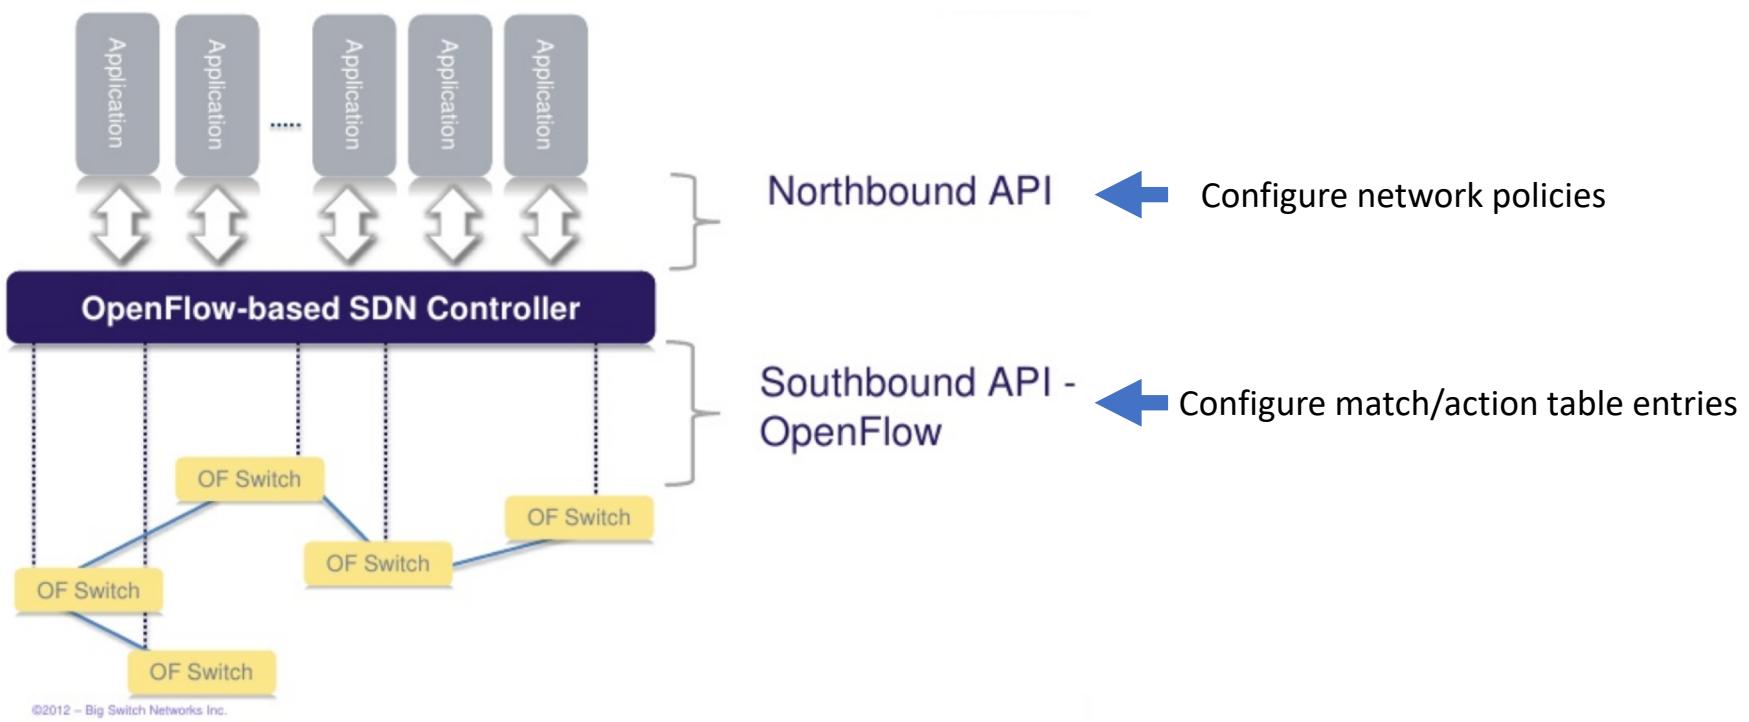
Switch Port	MAC src	MAC dst	Eth type	VLAN ID	IP Src	IP Dst	IP Prot	TCP sport	TCP dport	Action
*	*	00:1f..	*	vlan1	*	*	*	*	*	port6, port7, port9

# OpenFlow Example

Controller



# Some terminology



# Summarizing

- General principles
  - Controller can install **rules in switches**. Rules can match various **header fields**. A rule may entail **forwarding** packets to another switch/host, **dropping** packets, or **forwarding** packets to the controller
  - Typically, **any packet that does not match a rule is forwarded to the controller**

# Lots of flexibility!

- Fine vs. coarse grained rules
- Reactive vs. Proactive rule creation
- **There is no “default” or “best-practice” way to use OpenFlow!**

# Flow Routing vs. Aggregation

Both models are possible with OpenFlow

## Flow-Based

- Every flow is individually set up by controller
- Exact-match flow entries
- Flow table contains one entry per flow
- Good for fine grain control, e.g. campus networks

## Aggregated

- One flow entry covers large groups of flows
- Wildcard flow entries
- Flow table contains one entry per category of flows
- Good for large number of flows, e.g. backbone

# Reactive vs. Proactive

Both models are possible with OpenFlow

## Reactive

- First packet of flow triggers controller to insert flow entries
- Efficient use of flow table
- Every flow incurs small additional flow setup time
- If control connection lost, switch has limited utility

## Proactive

- Controller pre-populates flow table in switch
- Zero additional flow setup time
- Loss of control connection does not disrupt traffic
- Essentially requires aggregated (wildcard) rules

# OpenFlow Application: Network Slicing

- Divide the production network into **logical slices**
  - each slice/service controls **its own packet forwarding**
  - users pick **which slice controls their traffic**: opt-in
  - existing production services run in their own slice
    - e.g., Spanning tree, OSPF/BGP
- Enforce **strong isolation between slices**
  - actions in one slice do not affect another
- Allows the (logical) testbed to **mirror the production network**
  - real hardware, performance, topologies, scale, users

# Activity time!

- Form groups of 4/5 people
- Prepare a list of possible applications of OpenFlow (i.e., what can I do in an OpenFlow network that I could not do in a traditional network?)
  - You are encouraged to think beyond what discussed in the paper ☺

# OpenFlow beyond academia

# Academia & beyond

- OpenFlow was conceived as a research oriented solution
- ...but it has found application outside academia

# Commercial implementations

- OpenFlow is the first “mainstream” implementation of an idea called **software-defined networking (SDN)**
- Various SDN startups since 2007:
  - Nicira
  - Big Switch Networks
  - Embrane
  - ...
- **Takeway:** the ideas behind OpenFlow are relevant well beyond the academic world

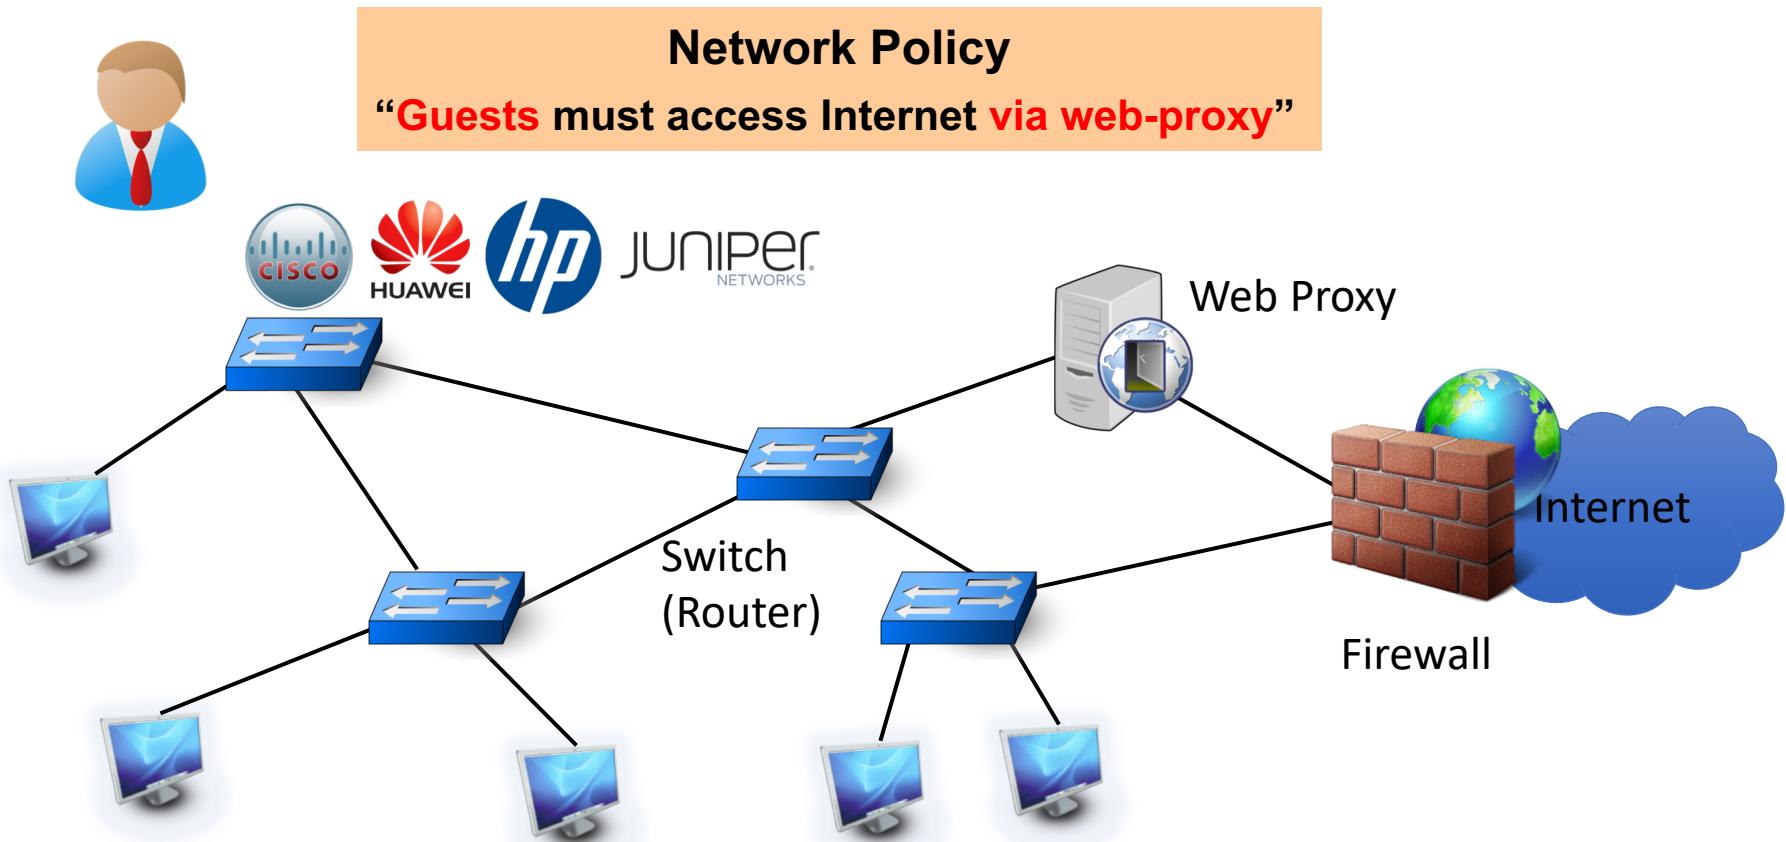
# So, what's the **real** point of SDN?

# SDN as a solution for network management

- SDN is relevant because it simplifies various tasks in **network management**

# Introduction

- Network management is driven by **policy** requirements



# Option1: Command-line Interface (CLI)

- Vendor-specific, low-level
- Script needs frequent maintenance



(Huawei)

```
Switch# display interface ethernet 1/0/0.1
Ethernet1/0/0.1 current state : UP
Line protocol current state : UP
Description : HUAWEI, AR Series, Ethernet1/0/0.1 Interface
.....
```

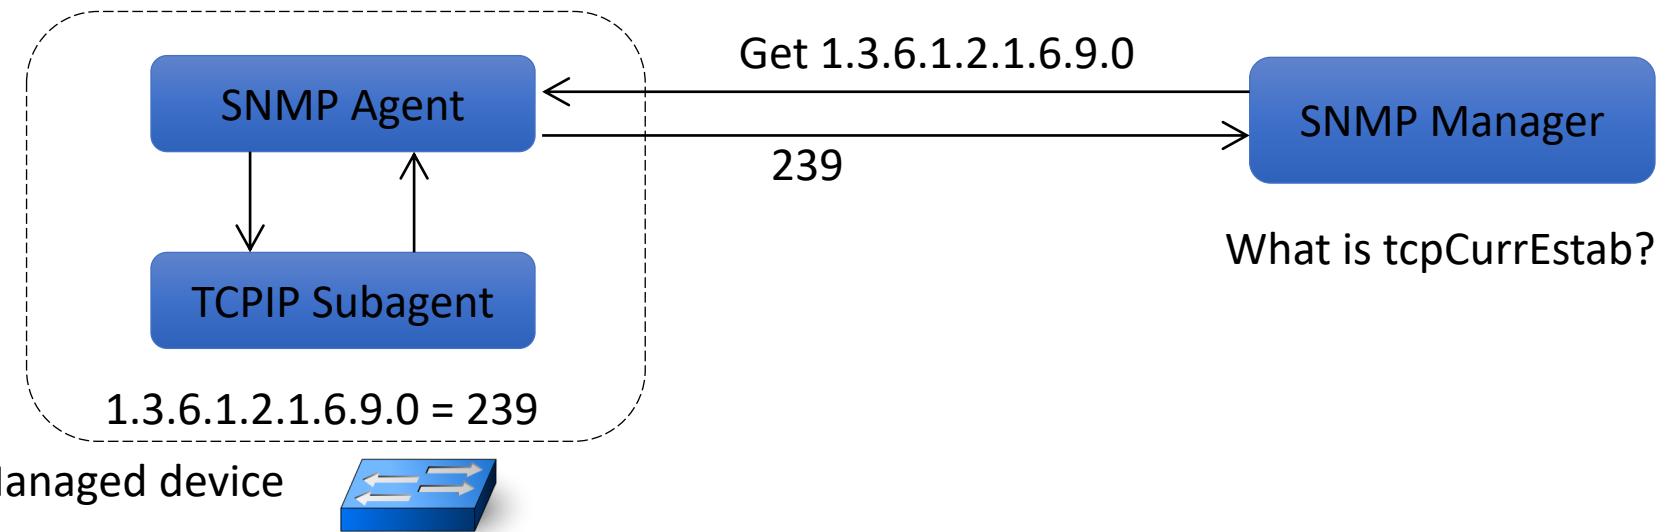
(Cisco)

```
Switch# show interfaces ethernet 0
Ethernet0 is up, line protocol is up
Hardware is Lance, address is 0060.3ef1.702b (bia 0060.3ef1.702b)
Internet address is 172.21.102.33/24
.....
```

**What could  
go wrong?**

# Option 2: Simple Network Management Protocol (SNMP)

- A framework that provides facilities for managing or monitoring network devices in IP networks
  - Query/Set provided variables in devices



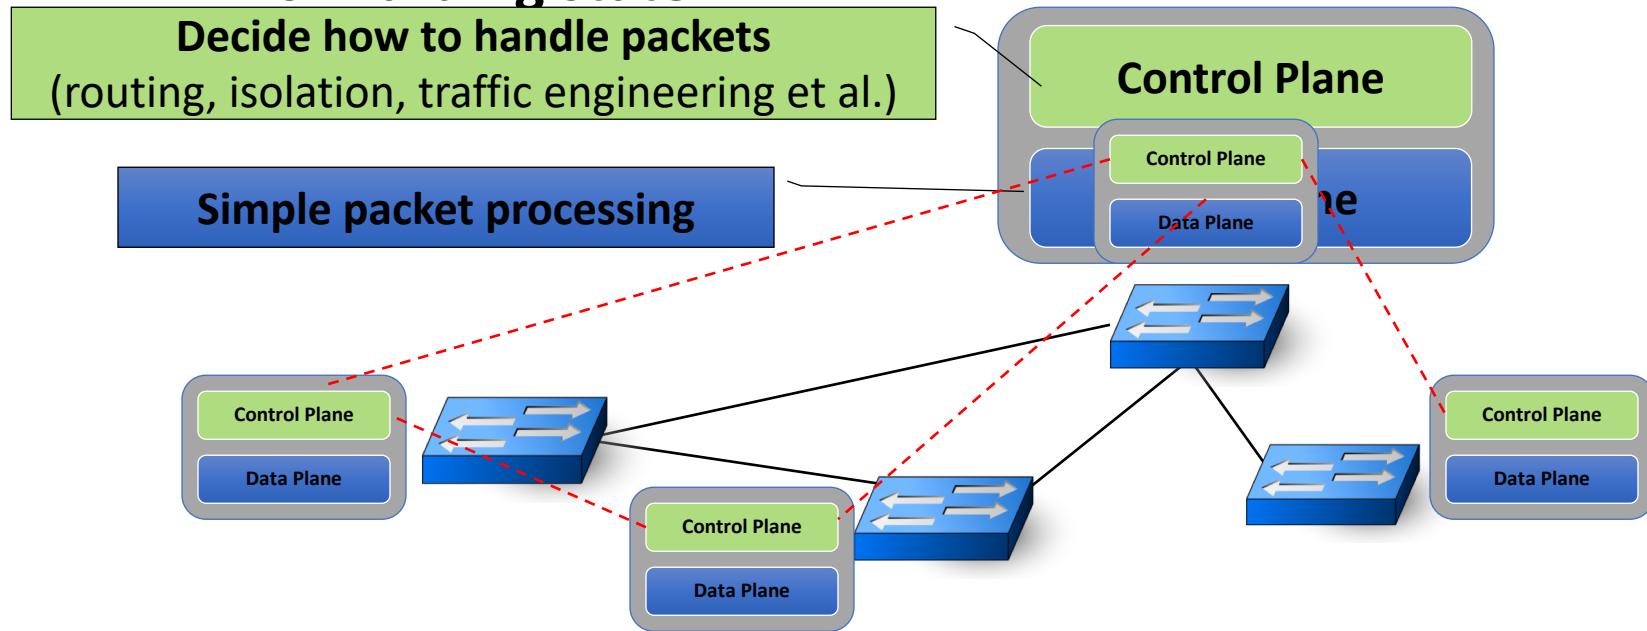
- Potential issues?

# Network Management is Challenging

- Current networks don't have good management approaches
  - CLI: vendor-specific, low-level 
  - SNMP: limited capabilities 
  - Third party management tools: bind to software vendor 
- Today's management requires
  - Easy configuration, fine-grained control, and react to continually changing network conditions timely
- Traditional approaches cannot satisfy today's management requirements

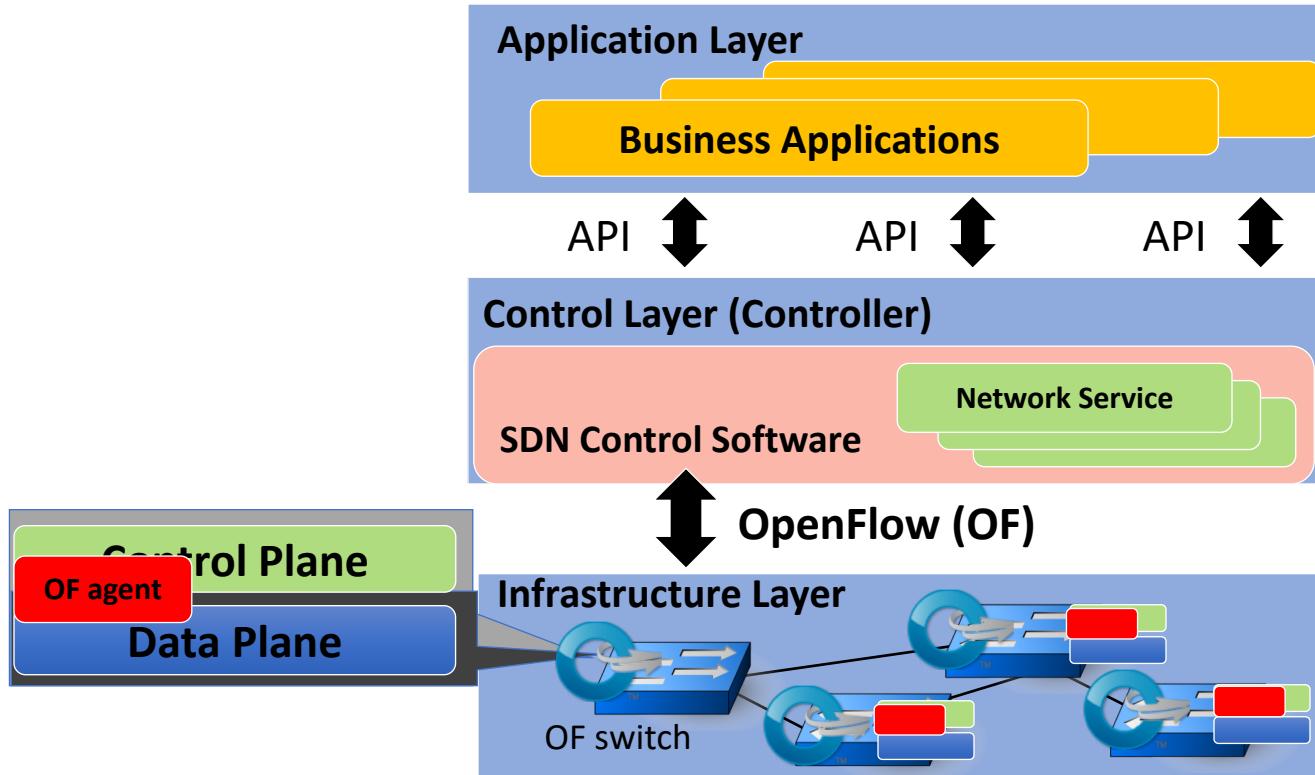
# What Is the Root Cause for the Complex Network Management?

- Control plane has **no good abstractions**
  - Okay for a single device, **not good for the whole network**
  - **No modularity and no standard way of defining forwarding state**

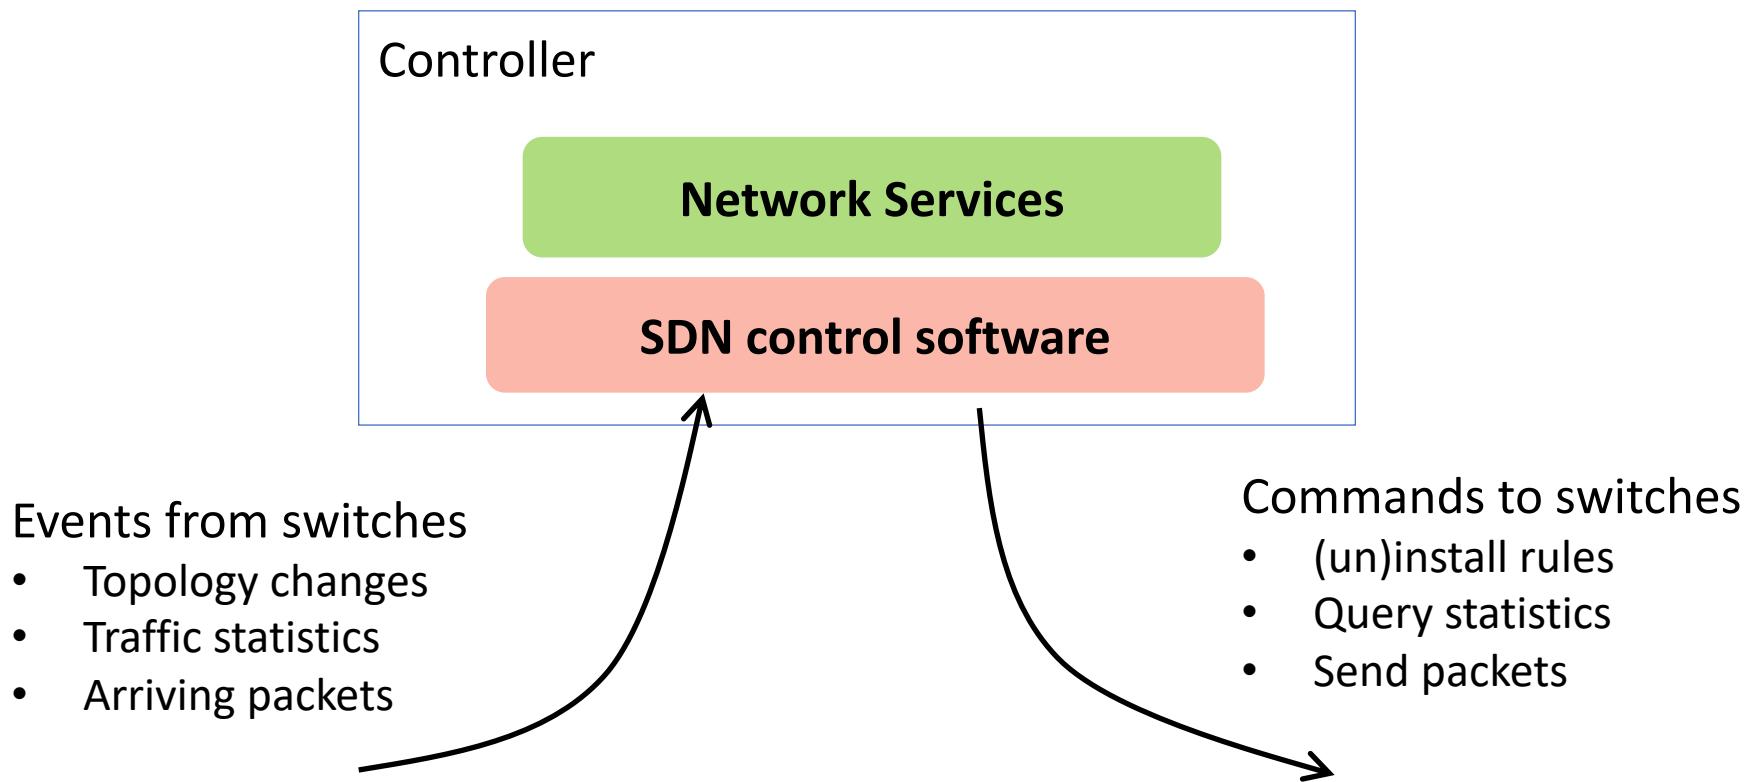


# Software Defined Networking (SDN) Architecture

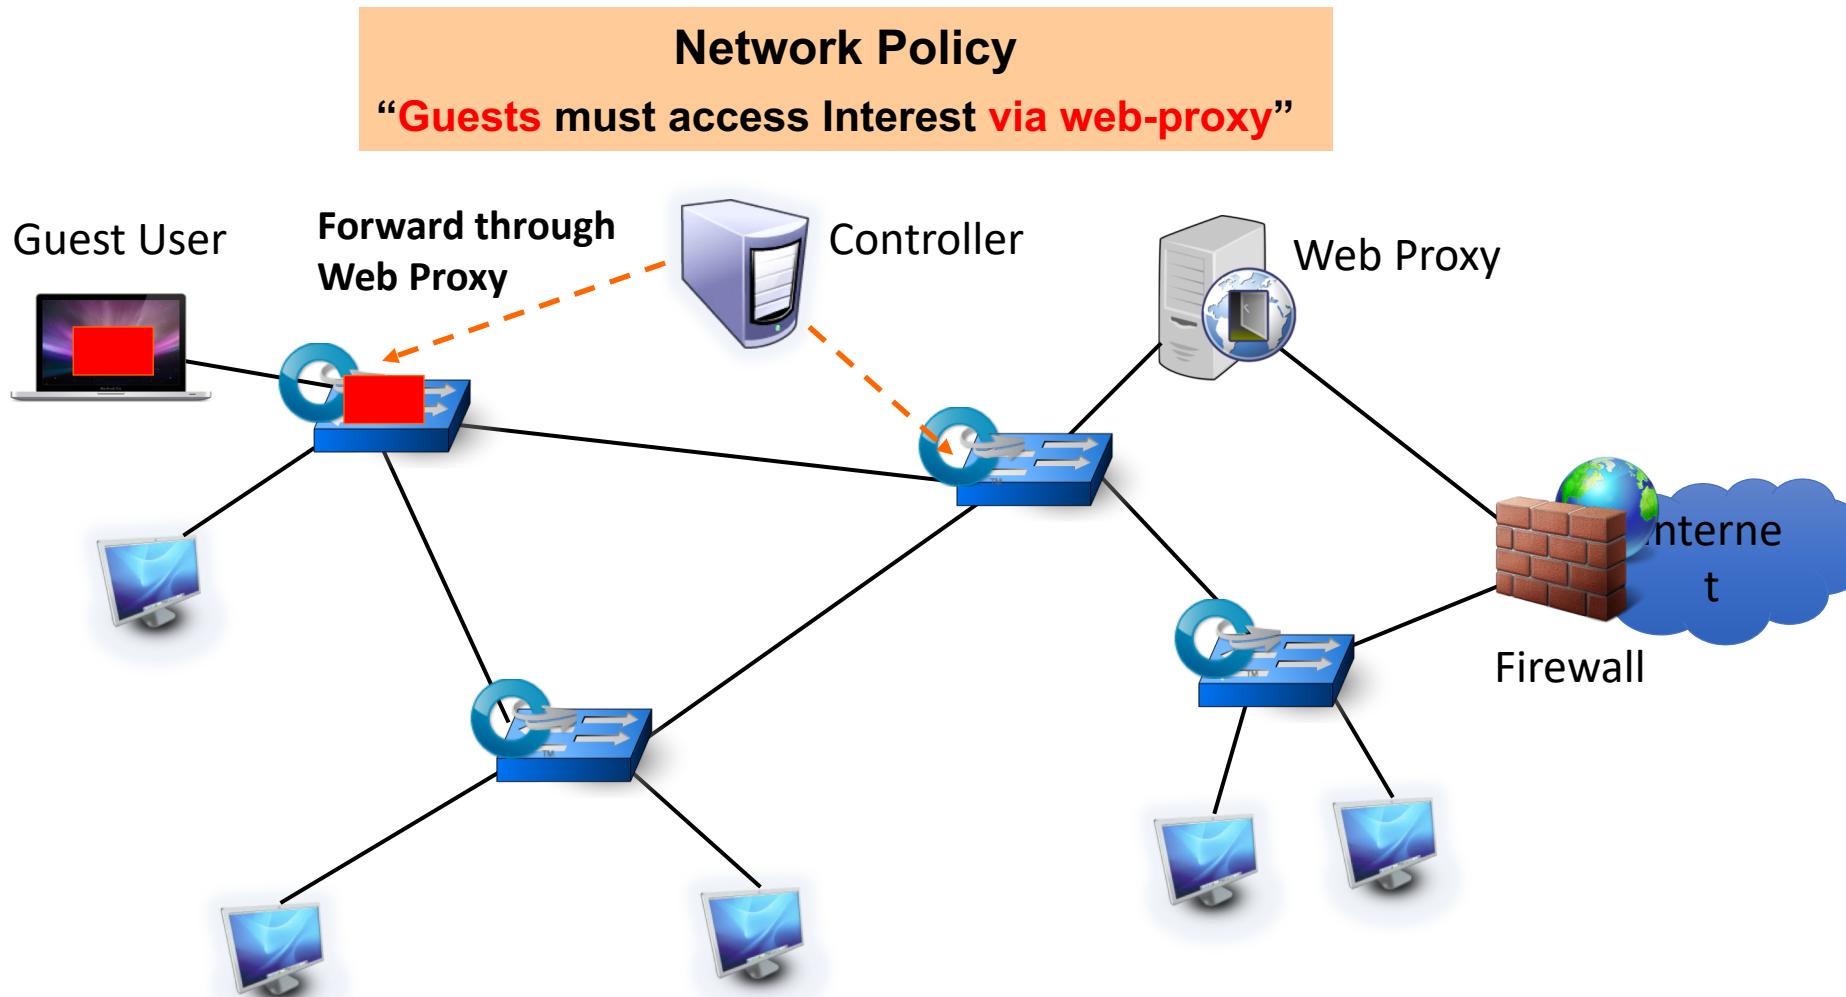
- A network architecture where network control is decoupled from forwarding and is directly programmable



# Programmable Control Plane



# Standard SDN Workflow



# Advantages of SDN

- Decouple control plane and data plane
  - The software control of the network can evolve independently of the hardware
- Logically centralized network control
  - Network-wide view, more deterministic, more efficient

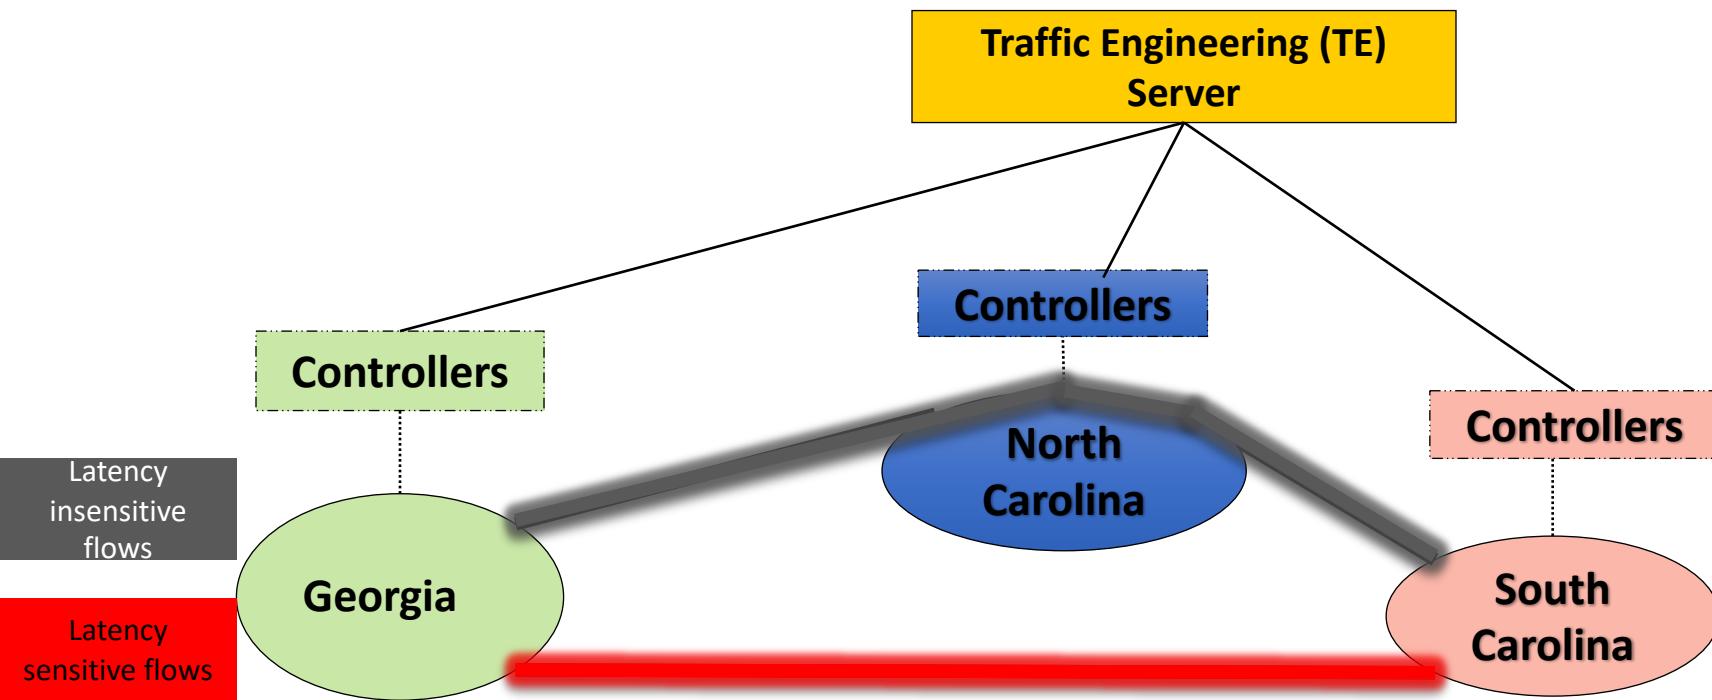
# A SDN Use Case

- Google inter-datacenters WAN (B4) is the first and largest SDN/OpenFlow deployment
- Overprovisioning introduces low link utilization (30-40%)



[3] B4: Experience with a globally-deployed software defined WAN. SigComm'13

# SDN Improves Link Utilization



- SDN drives link utilization to almost 100%
  - TE server assign flows to idle network links according to the application priority
  - Controllers enforce the assignment and control bursts

# Network Function Virtualization

- Take a physical device's actions and convert it to a VM performing the same tasks
  - Physical intrusion detection system becomes VM inspecting traffic
  - Physical firewall becomes VM with iptables
  - ...
- VM-based services can be scaled, deployed as needed etc.
- **SDN simplifies deployment of NFV (why?)**

# SDN Tools

- OpenFlow controllers:
  - Floodlight, OpenDayLight (Java)
  - Nox (C++)
  - Pox (Python)
  - ...
- Switches:
  - Open vSwitch (virtual software switch w/ OpenFlow support)
  - Hardware switches from various vendors

# Some limitations of software-defined networking

# Scalability Issues of SDN

- **Flow-level granularity places significant stress on switch state size and controller involvement**
- A switch can support a **limited number of OpenFlow entries**
  - OpenFlow rules are stored in TCAM (only support ~1500 entries in HP 5406zl switch)
  - Ethernet forwarding uses hash lookup in standard memory (support ~64k entries)
- Limited computation resources on controller
  - NOX controller can handle 30k flows/s
  - Datacenters with 100 edge-switches may have more than 10M flows/s

# Scalability Issues of SDN (cont.)

- Flow setup delay
  - Forward the initial packet: 5us on ASIC vs. 5ms if controller involved
  - **Unacceptable flow setup delay** for high performance networks
- **Frequent statistics-pulling** reduces the **flow setup rate**
  - Tasks (e.g. traffic engineering) need statistics gathering
  - 275 connections/s (no pull) vs. 50 connections/s (5 requests/s) using HP 5406zl switch

# Reliability issues of SDN

- **Let's hear some opinions first**
- Most significant issue: single-point-of-failure
  - The controller is responsible for managing the entire network
  - What if:
    - It cannot sustain the rate of requests from switches?
    - It is attacked?
    - ...or just crashes?

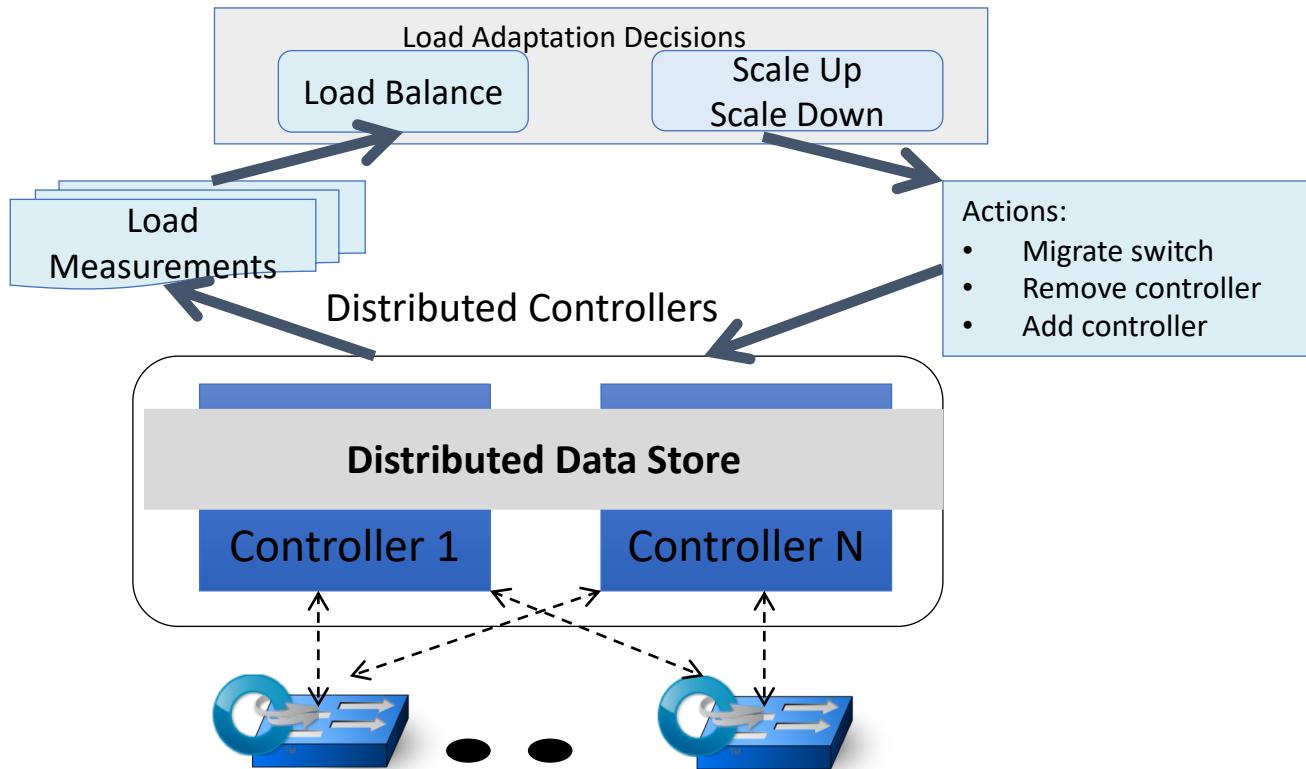
# Some possible solutions

# Scaling Control-plane

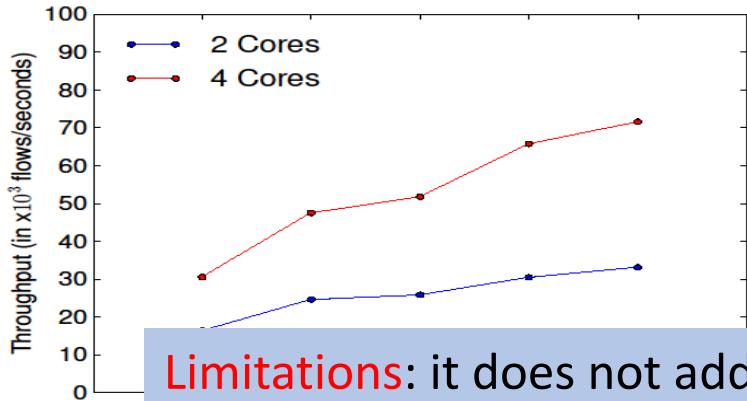
- Approaches in this category aim to **improve the limited computation resources on control layer**
  - Use better hardware, make it distributed, etc.
- **Distributed system** is better than centralized system
  - Better fault tolerance
  - Better scalability
  - Better load balancing

# Controller replication: ElastiCon

- **Make central controller a controller pool**
- Dynamically adds/removes controllers according to traffic conditions

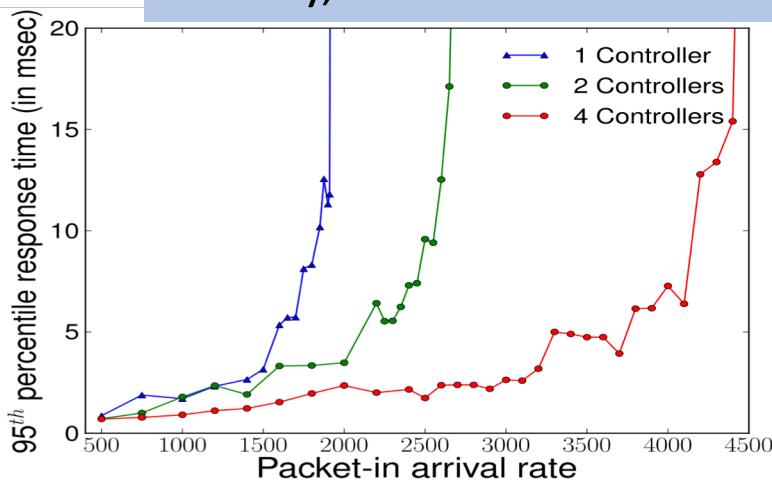


# Evaluation



Adding controller nodes increases the throughput

**Limitations:** it does not address the issue of flow setup delay directly, and introduces another layer of complexity



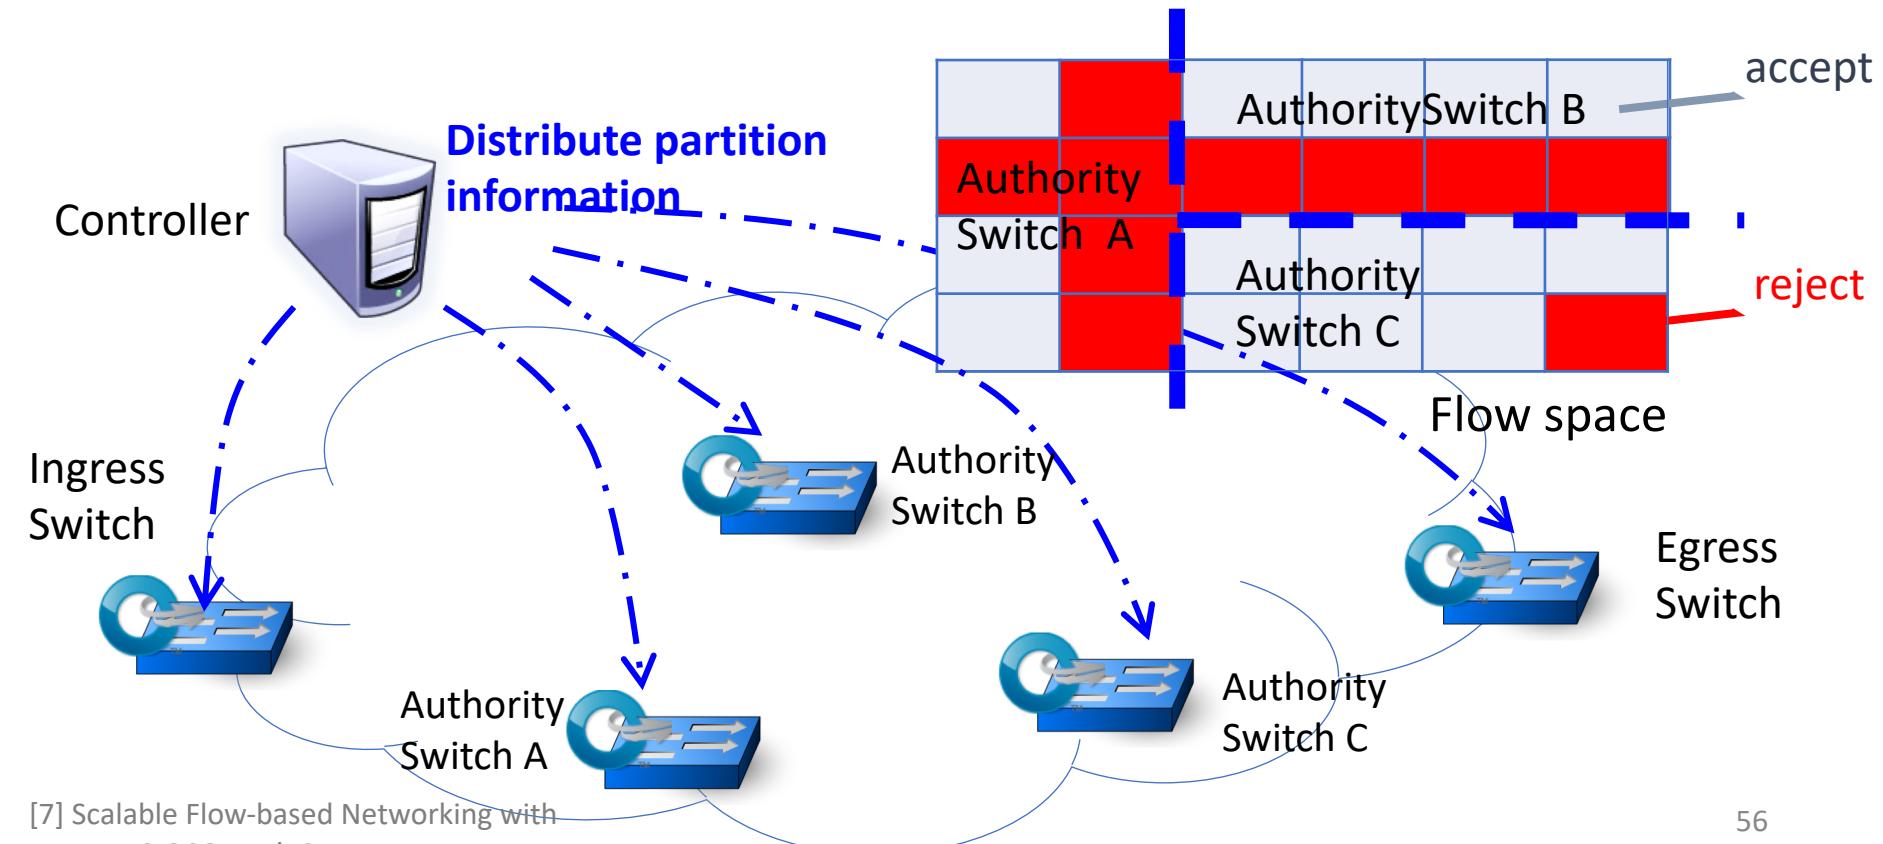
More controllers can make the response time stay in a certain level

# Making the controller less critical

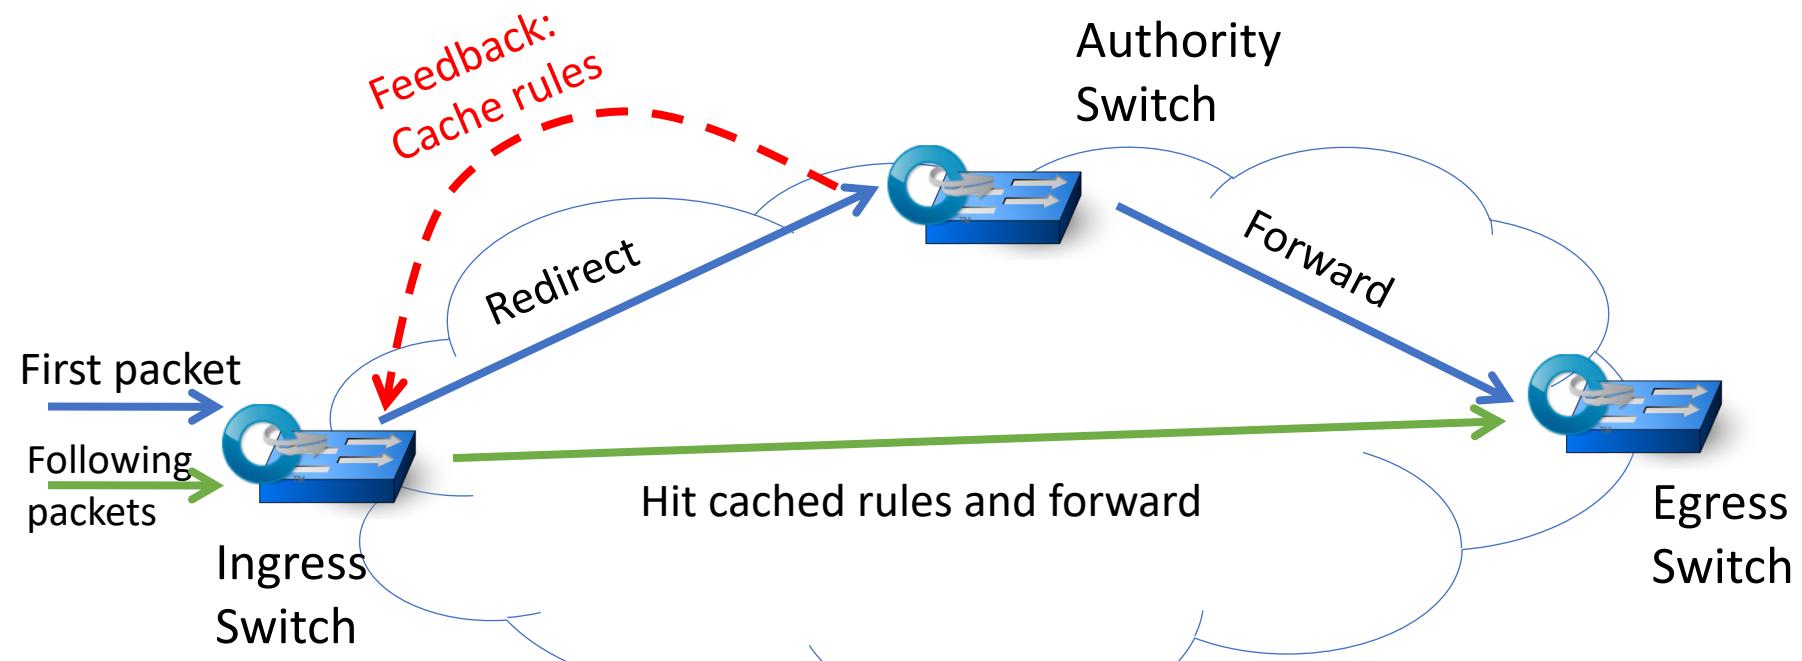
- **Delegation:** transfer some responsibility from controller to switches by making the switches “smarter”
- **Hybrid** between traditional network management and SDN

# DIFANE

- Controller delegates some switches to handle real-time jobs
  - Authority switches with larger TCAM space and processing capability
- Controller generates the rules, and installs them on switches

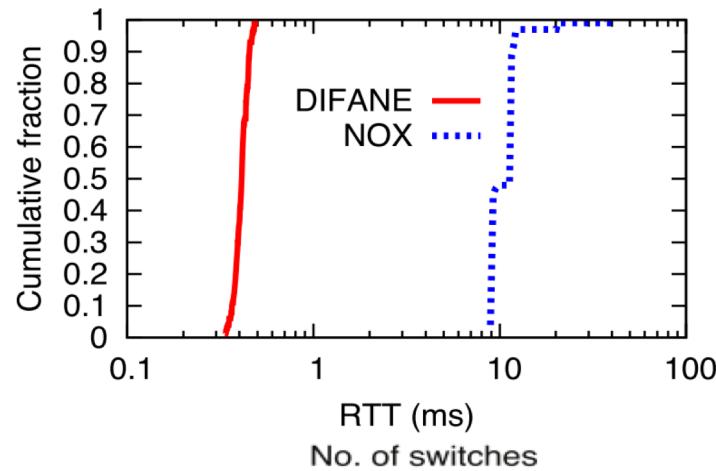


# DIFANE (cont.)

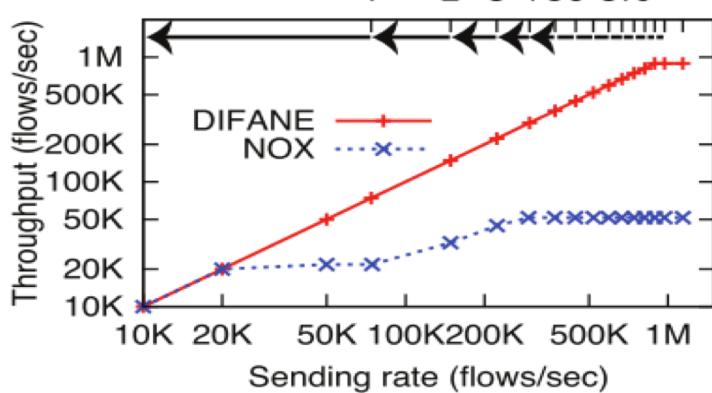


# DIFANE (cont.)

- DIFANE outperforms NOX ...



DIFANE achieves small delay for first packet (0.4ms vs 10ms)



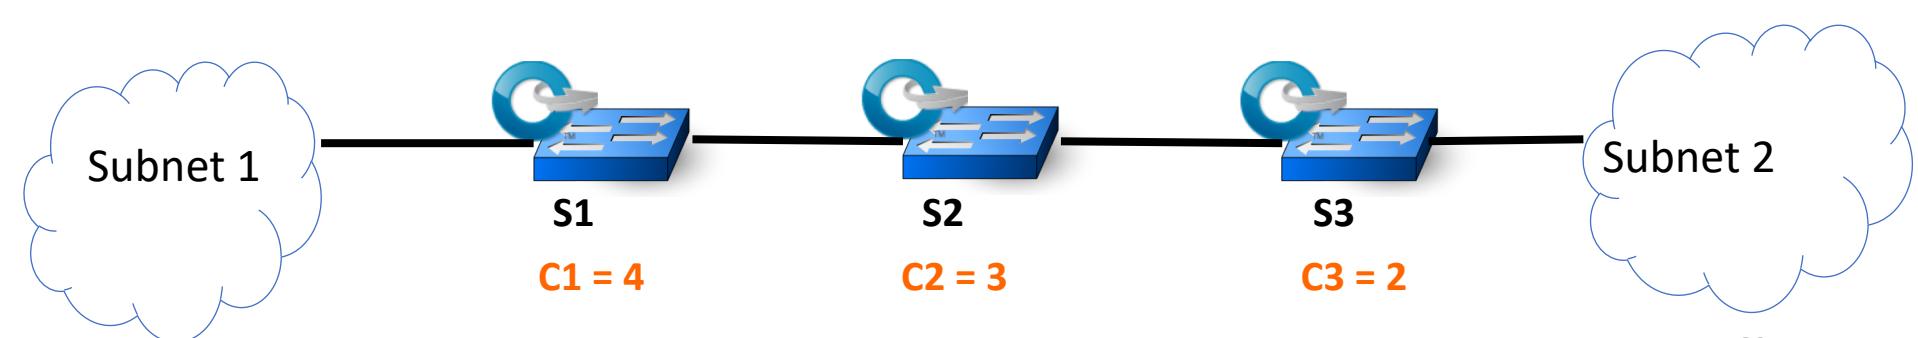
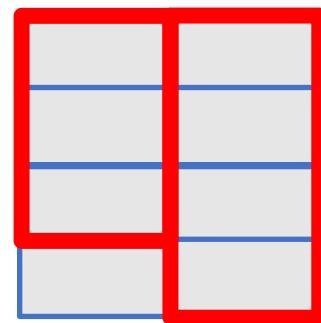
Throughput in DIFANE increases linearly with the number of switches.

# What about switch memory limitations?

- **Problem:** switch can only store a limited number of rules
  - What if the network policy requires more?
- **Possible solution:** only install rules on a per-needed basis
  - Generates a lot of work for controller
- **Alternate solution:** optimize rule placement
  - SDN controller naturally has the network wide view
  - Controller computes optimal placement of rules in switches' TCAMs
  - Advantages: less controller involvement, avoid TCAM space exhaustion
  - One big switch [8] abstraction

# One Big Switch

- Controller installs rules along the path while obeys the table size constraints
  - Assign subset of the endpoint rules (e.g. ACL) to a routing path
  - Place part of the rules on each switch along the path according to the table size



# One Big Switch (cont.)

- Synthetic 100-node topologies with real firewall configurations
- More than 60% unwanted traffic are dropped in the first 20% of the path
- **Limitations**
  - Introduce unwanted traffic
  - Statistic gathering is slow

# Conclusions

- Lots of hope when OpenFlow came out (~2008)
- Some of those hopes have materialized
  - See "success stories" like Google's
- **Lots of tricky issues**
- Vendors initially not very enthusiastic, then some support
  - Vendor-specific solutions integrating some SDN ideas
  - New buzzword: **policy-based networking**